

React Native Evals: A Reproducible Benchmark for Evaluating Large Language Models on Mobile App Development Tasks

Whitepaper

Artur Morys-Magiera, Mike Grabowski, Lech Kalinowski PhD,

Piotr Miłkowski, and Michał Pierzchała

Callstack, Poland

{artur.morys-magiera,mike,lech.kalinowski,
piotr.milkowski,michal.pierzchala}@callstack.com

April 17, 2026

Abstract

Large Language Models (LLMs) have improved rapidly on coding tasks, but most widely used benchmarks emphasize general-purpose programming and underrepresent mobile engineering. However, scores from generic coding benchmarks are often insufficient for estimating real-world performance on domain-specific tasks. To address this gap with respect to React Native, an open-source framework for building cross-platform applications, we present React Native Evals: a benchmark designed specifically for React Native problem solving evaluation. The framework evaluates declared requirements with optional weights, enabling fine-grained analysis of where models succeed or fail. These elements provide a transparent and auditable basis for comparative evaluation of LLMs on mobile app development tasks. The benchmark pipeline executes successfully end-to-end across the evaluated model set, confirming that the current eval suite is operationally usable. The results show that the framework can separate model capability in an auditable and reproducible way, while preserving requirement-level interpretability through explicit weights and per-eval verdict traces and deterministic output artifacts. Finally, we draw conclusions from the benchmark runs, comparing models both generally and within specific sub-domains in React Native code generation, as well as general metrics such as token consumption.

1 Introduction

Large Language Models (LLMs) have improved rapidly on coding tasks, but most widely used benchmarks emphasize general-purpose programming and underrepresent mobile engineering. React Native development introduces a different difficulty profile: correctness depends not only on code syntax, but also on component behavior, navigation flow, asynchronous state transitions, platform API usage, and consistency across multiple files. As a result, scores from generic coding benchmarks are often insufficient for estimating real-world performance on React Native tasks.

To address this gap, we present *React Native Evals*, a reproducible benchmark designed specifically for React Native code-generation evaluation. The benchmark uses repository-defined tasks, structured solver outputs, and requirement-level judging to assess whether generated implementations satisfy explicit engineering constraints. Rather than relying on a single pass/fail signal at the task level, the framework evaluates declared requirements with optional weights, enabling fine-grained analysis of where models succeed or fail.

This whitepaper documents the benchmark in an implementation-faithful manner relative to the current repository state. It specifies the dataset contract, end-to-end execution pipeline, solver and judge responsibilities, scoring and aggregation rules, artifact schema, and repeat-run protocol for variance-aware comparison. It also clarifies authoring conventions (including `implementation-requirement` IDs) and distinguishes those conventions from runner-enforced behavior. Together, these elements provide a transparent and auditable basis for comparative evaluation of LLMs on mobile app development tasks in React Native.

2 React Native methodology

The meritorious content of this benchmark is concentrated around validation of primitive Large Language Models without any external augmentations such as RAG, documentation access or agent skills. The purpose of this benchmark is to assess how those models perform in the task of React Native development including the most common third party libraries in the React Native ecosystem.

The benchmark we propose in this work includes evaluation cases (referred to as *evals*), divided into categories; each eval consists of multiple requirements as to the outputs of the evaluated model (referred to as *solver*), as described in sections 3.3 and 3.4. Finally, the solver performs the task in each evaluation, and in the second stage of the benchmark, another LLM (referred to as *judge*) assesses the result by assigning scores to each of the requirements, along with generating a justification in natural language (for interpretability). The breakdown of the categories of evals is presented in table 1. For the sake of this experiment, the chosen requirements were adjusted such that they mattered equally, thus implying weights of all requirements being set to 1.0.

Category	# Evals	Scope
Animation	14	React Native animation behavior, including <code>Animated</code> and <code>react-native-reanimated</code> .
Asynchronous state	14	Async state/data-flow tasks with <code>TanStack Query</code> , <code>Zustand</code> , and <code>Jotai</code> .
Lists	18	List rendering, item interactions, virtualization patterns, and sectioned data presentation.
Navigation	15	Screen transitions, stack/modal flows, route params, and navigation state handling.
React Native APIs	9	Core platform API usage such as layout, keyboard, accessibility, and device integration primitives.

Table 1: Benchmark categories, number of evals per category, and short scope descriptions.

3 The benchmark pipeline

This section specifies the exact models and runtime settings used in the benchmark to ensure reproducibility and fair comparison. We report model identifiers, version snapshots, assigned pipeline

roles, and key execution parameters. All experiments use the same evaluation protocol, task set, and scoring logic; only the configured model changes. The models used in this benchmark are presented in table 2.

3.1 Version and configuration of benchmarked models

Model name	Role
claude-opus-4.6	Solver
claude-opus-4.7	Solver
claude-sonnet-4.6	Solver, Judge
composer-2	Solver
composer-2-fast	Solver
deepseek-r1-distill-qwen-32b	Solver
deepseek-v3.2	Solver
gemini-3.1-pro-preview	Solver
gemma-4-31B-it	Solver
glm-5	Solver
gpt-5.3-codex	Solver
gpt-5.4	Solver
gpt-oss-20b	Solver
gpt-oss-120b	Solver
grok-4	Solver
kimi-k2.5	Solver
minimax-m2.7	Solver
qwen2.5-coder-32b-instruct	Solver

Table 2: List of models used in evaluation: names, versions and table role in the sense of section 3 (solver - the evaluated LLM; judge - LLM used to grade requirement execution by the solver)

3.2 Benchmark Unit and Dataset Contract

The atomic benchmark unit is a single eval directory. Each eval is expected to contain files structured according to the following sections 3.3 and 3.4. The general pipeline looks as follows:

1. **Generation CLI:** discover evals, load `app/` and `prompt.md`, run solver model, write generated files and `manifest.json` to the configured output directory.
2. **Judge CLI:** read eval entries from generation `manifest.json`, load generated files and `requirements.yaml`, run LLM judging, write per-eval results incrementally, and write summary at run completion.

Eval discovery is file-driven: the runner scans for `requirements.yaml`. If that file is missing, the eval is not discovered and will not be included in the run.

3.3 Generation stage inputs

1. **Discovery:** discover evals




Figure 1: End-to-end pipeline of the React Native Evals benchmark. Dataset inputs (blue) feed the Generation CLI (described in section 3.3), producing `manifest.json` and source files generated by the solver. The Judge CLI (described in section 3.4) reads the manifest and `requirements.yaml`, runs the judge LLM, maps verdicts by requirement ID, and writes `summary.json` along with telemetry data. Scoring computes metrics as described in section 6.

2. **Load `app/`:** read solution template (files on which the solver performs the task) and copy it to a temporary directory, in which the solver LLM will operate
3. **Load `prompt.md`:** parse the task definition, run solver model, write generated files and `manifest.json` to the configured output directory.

3.3.1 `prompt.md`

Task prompt used by the solver model. A representative logical template of this type of file is presented in listing 1.

`prompt.md` example

Implement X where performing interaction Y causes Z and state A uses an animation B from package C.

Listing 1: Example structure of a prompt in the benchmark.

3.3.2 `app/`

The `app/` directory must contain all files serving as baseline input provided to the solver. The files that shall be evaluated by the judge must be listed in the `requirements.yaml` file.

3.4 Judgement stage inputs

Reads eval entries from the generation output and runs an LLM to assess the completion of requirements (specified in section 3.4.2) in the generated solution.

3.4.1 `manifest.json`

The manifest file contains a summary of the generation stage (section 3.3), containing information according to listing 2.

3.4.2 `requirements.yaml`

Task prompt used by the solver model. A representative logical template of this type of file is presented in listing 3. The `version` is a constant element present for compatibility reasons for future versions of the project, that must be set to 1. The `inputs` list shall contain all files to be copied to the temporary workspace for the solver. The `requirements` list shall contain all the requirements, described in natural language. The requirements should explicitly state specific elements or behaviors that the presence of is to be asserted in the judgement stage. Optionally, each requirement may contain a `weight` value set to a positive value to emphasize its impact on the score for that requirement. If omitted, its value defaults to 1.

3.5 OpenCode integration in the judgment stage.

Execution layer used to call all solver and judge models through one interface (model IDs in the form `provider/model`, e.g., `openai/gpt-5.3-codex`). Before judging starts, the runner ensures that a local OpenCode server is available on the configured port; if a server is already running, it

manifest.json contract

```
{
  "runId": "YYYY-MM-DDTHH-MM-SS-SSSZ",
  "startedAt": "YYYY-MM-DDTHH:MM:SS.SSSZ",
  "finishedAt": "YYYY-MM-DDTHH:MM:SS.SSSZ",
  "solverModel": "<MODEL_NAME>",
  "pattern": "evals/**/*",
  "evalCount": 1,
  "evalsProcessed": 1,
  "evalsErrored": 0,
  "evals": [
    {
      "evalId": "01-category-task",
      "evalPath": "evals/category/task",
      "outputFiles": [
        "App.tsx",
        "styles.ts"
      ],
      "generatedPath": "category/task"
    }
  ]
}
```

Listing 2: The contract of `manifest.json` file, an artifact of the generation stage.

is reused. For each eval, the Judge CLI creates a new OpenCode session, sends the constructed judge prompt, and requests schema-constrained structured output. If strict structured output fails, a JSON fallback path is used and validated before scoring.

After each judge call, the pipeline can fetch OpenCode session metadata (conversation, token usage, and file-diff summary) and store it as `<eval-id>.opencode-session.judge.json`. Standard judge outputs are written incrementally as `runs/<input-folder>/evals/<eval-id>.json`, and the run-level aggregate is written as `summary.json`. Judging runs with bounded concurrency and retry logic; when `-fail-fast` is disabled, failed evals do not stop the run and are counted in `evalsErrored`.

The runner supports bounded concurrency. If an eval fails and `--fail-fast` is not enabled, the run continues and the eval is counted in `evalsErrored`.

The judge CLI also supports targeted re-judging over existing judge outputs using `--rerun-requirements-file` with `--output`; optionally pass `--rerun-requirement-id` to refresh only one requirement. Without `--rerun-requirement-id`, all requirements for the targeted eval are re-judged and replaced in the per-eval JSON. It also supports `--rerun-missing-judgements` with `--output` to scan for and judge all evals missing per-eval judge JSON outputs (the same missing rows counted by `evalsErrored` in the rebuilt summary). In rerun modes, the previous `summary.json` is backed up as `summary.backup.<timestamp>.json`, and a new aggregate summary is generated from the current per-eval result set.

```
requirements.yaml contract
```

```
version: 1
inputs:
  files:
    - app/App.tsx
    - app/styles.ts
requirements:
  - id: implementation-use-X-api
    description: Must use API `X` from package `A`.
    weight: 1
  - id: implementation-avoid-Y-api
    description: Must NOT use deprecated API `Y` from package `A`.
    weight: 1
  - id: toggle-state-on-click
    description: Upon click on button `B`, the visual active state of
      ↪ element `C` must toggle.
    weight: 1
```

Listing 3: The contract of `requirements.yaml` file.

4 Solver Stage Methodology

The solver methodology treats each benchmark item as a constrained code-generation task: the model receives a task prompt plus project context, produces a complete implementation artifact, and is evaluated only against task-defined requirements. The same solver protocol is applied across all tasks and runs, with outputs stored in a structured, reproducible format, ensuring fair comparison across models and direct traceability between generated code and downstream requirement-level scores.

4.1 Input Construction

The solver is given:

- the eval prompt from `prompt.md`
- the baseline files from `app/`
- a fixed system instruction requiring the model to return all provided files with modifications applied

The runner formats baseline files into a structured prompt with explicit file tags.

4.2 Output Contract

The solver returns structured output with:

- `summary`: short description of performed changes
- `files[]`: array of `{path, content}`

Returned paths are sanitized before writing to disk (e.g., stripping leading slashes and removing traversal segments such as `..`).

4.3 Model Requirement

The generation CLI requires `--model`. When set to `noop`, the generation stage copies files from each eval’s `reference/` directory into the configured output and still writes a validated `manifest.json`. Any other model value runs normal solver generation.

5 LLM Judge Methodology

For React Native coding tasks, an LLM-based judge is an effective approach that balances benefits with preparation workload. The quality of a task depends on semantic correctness across multiple files and layers. Component logic, state flow, navigation behavior, asynchronous handling, and API usage are all important. Many valid implementations exist, and strict deterministic rules do not apply to the semantic complexity of mobile apps. Deterministic checks are excellent for narrow invariants, but as a standalone judge, they overly penalize acceptable variations, miss errors in higher-level reasoning, and require constant revisions as frameworks and programming patterns evolve. A requirements- or specification-based LLM judge can evaluate requirements at the intent level, detect nuanced errors, incorrect cycle assumptions, or unstable query invalidation logic, and assign consistent, evidence-based results at the requirements level, thus remaining independent of the solver model and style. In practice, this approach provides a better fit to real-world engineering tasks and faster benchmark evolution than pure determinism pipelines for React Native.

5.1 Requirements Parsing

`requirements.yaml` is parsed and validated at runtime. The enforced schema requires:

- **version**: positive integer (defaults to 1 via schema default)
- **requirements[]**: non-empty list of objects with:
 - **id** (non-empty string)
 - **description** (non-empty string)
 - optional positive **weight**

Extra fields are ignored by the current runtime parser.

5.2 Deterministic implementation- Requirement Convention

The repository documents an authoring convention for deterministic requirements:

- prefix deterministic, source-verifiable requirement IDs with **implementation-**
- use **weight: 1** by default to avoid overweighting implementation details

Examples include required imports, API usage patterns, and specific wiring that can be verified directly from files.

Important distinction: the runner does *not* infer behavior from ID prefixes such as **implementation-**. Prefixes are an authoring convention; scoring changes only when a requirement explicitly sets **weight** in `requirements.yaml`.

5.3 Technical Eval Authoring Controls

For technical categories (for example navigation, animation, storage), requirement quality is governed by additional authoring controls documented under `docs/`.

- **Recent API shift audit:** before requirement writing, review official docs plus recent releases/changelogs for each core library, and extract concrete “new preferred API” and “deprecated/removed API” signals with dates.
- **Evidence-gated MUST NOT:** use explicit MUST NOT clauses only when backed by primary-source deprecation/removal or correctness caveats.
- **Implementation-level specificity:** technical diversity is measured at requirement/API-constraint level, not prompt wording level.
- **Similarity budget:** within a library subgroup, keep shared baseline requirements small (target: at most two shared IDs) and maintain multiple eval-specific implementation constraints per eval.
- **Judge input scope:** set `inputs.files` to implementation files under evaluation (for example `app/App.tsx`) rather than prompt-only artifacts.

These controls are authoring policy for dataset quality; they are not independently enforced by runner schema validation.

5.4 Judge Prompt Construction

The judge prompt includes:

- declared requirements (including weights when present)
- generated files under evaluation
- rules instructing the judge to use only provided files as evidence and return one result per declared requirement ID

5.5 Structured Judge Output

The judge returns structured output with:

- optional `summary`
- `requirements[]`, where each row includes `id`, `passed`, `reason`, `evidence[]`, and optional `confidence`

5.6 Requirement Mapping and Failure Policy

Judge rows are mapped back to declared requirements by ID.

- If a declared requirement is missing from judge output, the runner marks it as failed.
- The runner emits a fixed failure reason indicating that the judge did not return a result for that requirement.
- Extra judge rows for undeclared IDs are ignored for scoring.
- Requirement weights are normalized before scoring (section 6).

The judge CLI requires `--model`; there is no judge skip/noop path.

6 Scoring Methodology

Each model is evaluated through the identical benchmark pipeline in 10 independent runs to reduce variance from stochastic generation and judging effects and to produce a statistically more reliable performance estimate. In every run, the model receives the same task set, prompt protocol, and scoring rules, while random factors may still lead to different outputs and per-task scores. Repeating the full pipeline therefore captures both central tendency and dispersion of performance, allowing analysis to rely on aggregated results (e.g., mean and spread) rather than a potentially noisy single-run outcome. This design improves fairness in cross-model comparison and supports more robust conclusions about relative capability on React Native tasks.

6.1 Requirement Weight Normalization

Each requirement weight is normalized as:

$$w'_i = \begin{cases} w_i & \text{if } w_i \text{ is finite and } w_i > 0 \\ 1 & \text{otherwise} \end{cases}$$

This normalization is defensive. In normal operation, the runtime schema already requires positive weights when provided.

6.2 Per-Eval Weighted Requirement Score

For an eval with declared requirements R , normalized weights w'_i , and binary pass indicators p_i :

$$\begin{aligned} \text{totalWeight} &= \sum_{i \in R} w'_i \\ \text{passedWeight} &= \sum_{i \in R} p_i \cdot w'_i \\ \text{scoreRatio} &= \begin{cases} 0 & \text{if totalWeight} = 0 \\ \frac{\text{passedWeight}}{\text{totalWeight}} & \text{otherwise} \end{cases} \end{aligned}$$

The runner rounds `passedWeight`, `totalWeight`, and `scoreRatio` to 4 decimal places.

Interpretation example. If requirements have weights $[2, 2, 1]$ and pass outcomes $[1, 0, 1]$, then the per-eval score is:

$$\frac{2 + 0 + 1}{2 + 2 + 1} = \frac{3}{5} = 0.6$$

This pattern is relevant only when authors intentionally assign non-default weights to specific requirements.

6.3 Run-Level Aggregate Score

Let E^+ be the set of evals that complete successfully (i.e., without an unrecovered pipeline error). The summary field `weightedAverageScore` is:

$$\text{weightedAverageScore} = \begin{cases} 0 & \text{if } |E^+| = 0 \\ \frac{1}{|E^+|} \sum_{e \in E^+} \text{scoreRatio}_e & \text{otherwise} \end{cases}$$

rounded to 4 decimal places.

Interpretation note: the “weighted” part refers to weights within each eval. The run-level average itself is an unweighted mean across successful evals.

6.4 Additional Summary Metrics

The run summary also reports:

- `requirementsTotal`
- `requirementsPassed`
- `evalsErrored`

For successful evals, each row includes `requirementsTotal`, `requirementsPassed`, and `scoreRatio`. For pipeline errors, each row includes `status=error` along with the eval identifier and path.

7 Artifacts and Output Schema

Each run writes artifacts to deterministic generation and judge directories:

- `runs/<input-folder>/evals/<eval-id>.json`: per-eval judge results (written as each eval completes)
- `runs/<input-folder>/summary.json`: aggregate judge summary (written at end)
- generated artifacts and `manifest.json`: written by generation CLI to a user-selected output directory (default: `generated/<model>-<run-id>`)
- per-eval OpenCode session snapshots:
 - generation: `<generated-eval-dir>/opencode-session.solver.json`
 - judging: `runs/<input-folder>/evals/<eval-id>.opencode-session.judge.json`
- optional debug artifacts when `--debug` is enabled

Per-eval results include solver and judge model identifiers, mapped requirement verdicts, per-eval score, generated file list, and optional `judgeSessionArtifactPath` when session capture succeeds. The generation `manifest.json` includes `solverSessionArtifactPath` for each processed eval.

8 Repeat Orchestration Script

The repository provides `scripts/bench-series.sh` to run multiple full cycles (generation + judging) for variance analysis. Configure via environment variables:

- `RUN_COUNT`: number of run+judge iterations (default: 10)
- `BENCH_RUN_ARGS`: arguments for the generation CLI (e.g. `--model`, `--pattern`)
- `BENCH_JUDGE_ARGS`: arguments for the judge CLI (e.g. `--model`); `--input` is injected per run
- `BENCH_OUTPUT_BASE`: base directory for outputs (default: `/tmp/bench-series`)

For each of the `RUN_COUNT` iterations, the script runs generation then judging. If either step fails, it retries up to 3 times before failing the script. Outputs are written to `<BENCH_OUTPUT_BASE>/<timestamp>/run-1`, `run-2`, etc.

9 Results

Figure 2 summarizes both overall and category-level performance across the evaluated solver models. In overall weighted average score, **composer-2** leads (96.1%), followed by **composer-2-fast** (94.9%) and **gpt-5.4** (85.3%), while **deepseek-r1-distill-qwen-32b** scores lowest (44.4%). The category breakdown remains non-uniform: navigation and React Native API tasks are generally the strongest categories across the model set, whereas animation remains the most discriminative category and lists provide additional separation among mid-tier models. For each model, detailed analysis is provided in appendix section A. The table 4 summarizes token usage for each of the models.




Figure 2: Overall score and category comparison across all models.

Rank	Model	Weighted average score	Runs
1	composer-2	0.961 (96.1%)	10
2	composer-2-fast	0.949 (94.9%)	10
3	gpt-5.4	0.853 (85.3%)	10
4	claude-opus-4.6	0.841 (84.1%)	10
5	gpt-5.3-codex	0.831 (83.1%)	10
6	claude-opus-4.7	0.828 (82.8%)	10
7	claude-sonnet-4.6	0.806 (80.6%)	10
8	gemini-3.1-pro-preview	0.789 (78.9%)	10
9	kimi-k2.5	0.772 (77.2%)	10
10	gemma-4-31B-it	0.752 (75.2%)	10
11	glm-5	0.748 (74.8%)	10
12	grok-4	0.726 (72.6%)	10
13	gpt-oss-120b	0.716 (71.6%)	10
14	deepseek-v3.2	0.715 (71.5%)	10
15	minimax-m2.7	0.714 (71.4%)	10
16	gpt-oss-20b	0.710 (71.0%)	10
17	qwen2.5-coder-32b-instruct	0.512 (51.2%)	10
18	deepseek-r1-distill-qwen-32b	0.444 (44.4%)	10

Table 3: Overall model comparison from packed result archives.

Model	Tokens mean	Tokens stddev
claude-opus-4.6	885,374	1,537
claude-opus-4.7	1,209,561	723
claude-sonnet-4.6	885,424	1090
composer-2	—	—
composer-2-fast	—	—
deepseek-r1-distill-qwen-32b	254,437	1,247
deepseek-v3.2	5,590,132	782,834
gemini-3.1-pro-preview	1,074,284	12,826
gemma-4-31B-it	885,910	91,050
glm-5	1,171,183	152,778
gpt-5.3-codex	778,463	49,875
gpt-5.4	869,589	50,085
gpt-oss-120b	653,327	27,343
gpt-oss-20b	765,263	57,402
grok-4	1,257,083	38,612
kimi-k2.5	2,500,716	408,562
minimax-m2.7	1,261,454	131,780
qwen2.5-coder-32b-instruct	732,531	15,154

Table 4: Total tokens per run by model. Mean and stddev computed across runs. It was impossible to **precisely** determine Composer 2 and Composer 2 Fast values for the runs.

The full benchmark pipeline executed successfully end-to-end across the evaluated model set, confirming that the current eval suite is operationally usable for comparative React Native model assessment.

10 Recommended Reporting Protocol

For comparative studies, report at minimum:

1. repository commit hash (dataset and runner version)
2. CLI options (run: `--pattern`, `--model`, `--timeout`, `--concurrency`, `--output`;
judge: `--model`, `--timeout`, `--concurrency`, `--input`
and optional rerun flags `--rerun-requirements-file`, `--output`,
scope optional `--rerun-requirement-id`, scope optional `--rerun-missing-judgements`)
3. execution date and time
4. counts of discovered, processed, and errored evals
5. `weightedAverageScore` and `requirementsPassed/Total`
6. generated artifact source directory

Repeat runs are recommended. The `bench:repeat` script (or `scripts/bench-series.sh`) automates multiple cycles. The current runner does not explicitly pin a random seed or decoding settings (e.g., temperature) in solver or judge calls, so provider defaults and model nondeterminism can affect results.

11 Limitations and Threats to Validity

- **Judge dependence:** requirement verdicts depend on the configured LLM judge model and prompt behavior.
- **Requirement-interpretation ambiguity:** without seeded reference context, ambiguous requirement phrasing can increase variance in judge decisions.
- **Binary requirement outcomes:** there is no partial credit within a single requirement.
- **Run-level averaging choice:** the summary score is unweighted across evals, which can underweight larger evals.
- **Pipeline error exclusion from mean:** errored evals are counted separately and excluded from the run-level score average.
- **Authoring-convention variability:** conventions such as `implementation` improve consistency but are not enforced by the runtime parser.

12 Acknowledgments

Text drafting and editorial refinement used artificial intelligence-assisted tooling. All technical claims, implementation details, equations, and reported metrics were manually verified against repository code and captured experiment artifacts before inclusion.

13 Conflict of Interest

The authors are affiliated with Callstack. The experiments used commercially available infrastructure acquired and operated independently; no direct sponsorship, funding, or editorial influence was provided for this work.

14 Conclusion

We evaluated **18** solver models on the React Native eval suite (**5** task categories in total), with **10 repeated runs per model** under the same pipeline and scoring rules. Across completed runs, the highest mean `weightedAverageScore` was **0.961 (96.1%)**, model: `composer-2`; the median across models was **0.761 (76.1%)**. Category-level analysis (fig. 2) shows strongest performance on **navigation** and **react native apis**, while **animation** remains the weakest and most discriminative category; **lists** and **async state** provide additional separation among otherwise closely clustered models. Recurring failures were concentrated in behavior-heavy requirements, especially animation dynamics and async-state consistency/invalidation logic.

Two practical implications follow. First, React Native development difficulty is not uniform: models that are competitive on navigation and structural code generation can still be unreliable on behavior-heavy work, where correctness depends on timing, gesture/animation dynamics, and state invalidation semantics. For practitioners, this suggests applying the most scrutiny (and the most automated protection via tests and runtime validation) to animation and async-state code paths, even when a model performs strongly overall. Second, requirement-level evaluation yields actionable signals: per-requirement verdict traces support targeted prompt iteration, scaffolding improvements, and dataset evolution without collapsing all behavior into a single task-level pass/-fail.

From an operational standpoint, token usage varied substantially across models (table 4). Among models with available token aggregates, mean usage ranged from **254,437** tokens per run (`deepseek-r1-distill-qwen-32b`) to **5,590,132** (`deepseek-v3.2`), with additional high-consumption behavior visible for `kimi-k2.5` at **2,500,716**. The token totals for `composer-2` and `composer-2-fast` could not be determined precisely from the captured runs, so efficiency comparisons for those two models should be interpreted with caution. Overall, this highlights a cost/latency frontier that is not directly implied by score alone, and reinforces that model selection for production assistance should consider both category-specific performance and consumption characteristics.

Overall, these results indicate that React Native Evals separates model capability in a reproducible and auditable way while preserving requirement-level interpretability (via explicit per-requirement weights and per-eval verdict artifacts). They also highlight a need to extend the current evaluation set beyond the existing categories to more of the day-to-day React Native surface area, such as handling network requests (including error handling, retries, caching, and cancellation), building custom native integrations, repairing or extending faulty code in partially-correct code-bases, keeping pace with the latest versions and breaking changes of common third-party libraries, and more. Future work should also explore judge-robust scoring (e.g., multi-judge consensus and calibrated partial credit) to reduce dependence on any single evaluator while maintaining transparency.

A Detailed Results

A.1 claude-opus-4.6

Model	Runs	Score	Reqs	P@1	P@5	P@10
vercel/anthropic/claude-opus-4.6	10	84.1%	2296/2730	60.6%	66.8%	69.7%

Table 5: Model summary for claude-opus-4.6

Category	Score	Evals	Requirements	Failed requirements
animation	77.8%	130	53	121
async state	83.6%	130	65	106
lists	80.8%	180	79	160
navigation	94.5%	130	45	28
react native apis	93.0%	90	31	19

Table 6: Results by category for claude-opus-4.6. Score: mean(per-eval passed/total weight) across evals and runs. Evals: count of evaluation tasks. Requirements: sum of requirement count per eval. Failed requirements: sum of failed requirement instances across runs.

Eval	Category	Score	P@1	P@5	P@10
01-rn-anim-pressable-scale-with-timing	animation	100.0%	100.0%	100.0%	100.0%
02-rn-anim-spring-toggle-card-height	animation	83.3%	50.0%	99.6%	100.0%
03-rn-anim-layout-enter-exit-list-items	animation	100.0%	100.0%	100.0%	100.0%
04-rn-anim-scroll-linked-collapsing-header	animation	100.0%	100.0%	100.0%	100.0%
05-rn-anim-pan-drag-card-with-snap-points	animation	77.5%	10.0%	50.0%	100.0%
06-rn-anim-longpress-then-pan-activation-gate	animation	100.0%	100.0%	100.0%	100.0%
07-rn-anim-pinch-pan-simultaneous-photo-canvas	animation	100.0%	100.0%	100.0%	100.0%
09-rn-anim-android-modal-gesture-root-safety	animation	83.3%	0.0%	0.0%	0.0%
10-rn-anim-worklet-thread-bridge-schedule-on-rn	animation	37.5%	0.0%	0.0%	0.0%
11-rn-rngh-exclusive-single-double-tap	animation	62.0%	0.0%	0.0%	0.0%
12-rn-rngh-scroll-pan-failure-chain	animation	80.0%	0.0%	0.0%	0.0%
13-rn-worklets-runtime-offload-and-bridge	animation	33.3%	0.0%	0.0%	0.0%
14-rn-keyboard-controller-aware-scroll-footer	animation	54.0%	0.0%	0.0%	0.0%
01-rn-rq-stable-query-key-filters	async state	56.0%	0.0%	0.0%	0.0%
02-rn-rq-dependent-query-enabled-gate	async state	58.0%	0.0%	0.0%	0.0%
03-rn-rq-mutation-invalidate-on-success	async state	44.0%	0.0%	0.0%	0.0%
04-rn-rq-optimistic-update-rollback	async state	95.7%	70.0%	100.0%	100.0%
05-rn-zustand-async-action-lifecycle	async state	100.0%	100.0%	100.0%	100.0%
06-rn-zustand-request-id-stale-response-guard	async state	76.0%	0.0%	0.0%	0.0%
07-rn-zustand-persist-async-hydration-gate	async state	88.6%	30.0%	91.7%	100.0%
09-rn-jotai-async-read-atom-suspense	async state	100.0%	100.0%	100.0%	100.0%
10-rn-jotai-loadable-async-atom-no-suspense	async state	72.5%	20.0%	77.8%	100.0%
11-rn-jotai-async-write-atom-mutation	async state	100.0%	100.0%	100.0%	100.0%
12-rn-jotai-atomwithstorage-async-rehydrate-guard	async state	100.0%	100.0%	100.0%	100.0%
13-rn-react-usedeferredvalue-search-stale-indicator	async state	100.0%	100.0%	100.0%	100.0%
14-rn-rq-react-native-appstate-focus-online-managers	async state	96.7%	90.0%	100.0%	100.0%
01-rn-flatlist-scroll-to-unread-fixed-rows	lists	100.0%	100.0%	100.0%	100.0%
02-rn-flatlist-basic-contact-list	lists	100.0%	100.0%	100.0%	100.0%
03-rn-flatlist-two-column-product-grid	lists	100.0%	100.0%	100.0%	100.0%
04-rn-flatlist-viewability-impression-tracking	lists	80.0%	0.0%	0.0%	0.0%
05-rn-flatlist-large-feed-window-tuning	lists	100.0%	100.0%	100.0%	100.0%
06-rn-flatlist-pull-to-refresh-feed	lists	82.0%	10.0%	50.0%	100.0%
07-rn-flatlist-to-sectionlist-conversion	lists	100.0%	100.0%	100.0%	100.0%
08-rn-flashlist-basic-list	lists	100.0%	100.0%	100.0%	100.0%
09-rn-flashlist-horizontal-rail-list-header	lists	100.0%	100.0%	100.0%	100.0%
10-rn-flashlist-horizontal-dynamic-height-header-first-item	lists	50.0%	0.0%	0.0%	0.0%
11-rn-flashlist-expandable-cards-layout-state	lists	66.7%	0.0%	0.0%	0.0%
12-rn-flashlist-chat-start-from-bottom	lists	40.0%	0.0%	0.0%	0.0%
13-rn-flashlist-masonry-photo-gallery	lists	40.0%	0.0%	0.0%	0.0%
14-rn-flatlist-to-legendlist-migration	lists	70.0%	10.0%	50.0%	100.0%
15-rn-legendlist-sticky-grouped-agenda	lists	80.0%	0.0%	0.0%	0.0%
16-rn-legendlist-section-scroll-to-location	lists	100.0%	100.0%	100.0%	100.0%
17-rn-legendlist-chat-start-from-bottom	lists	96.0%	80.0%	100.0%	100.0%
18-rn-legendlist-expandable-cards-recycling-state	lists	50.0%	0.0%	0.0%	0.0%
01-rn-nav-stack-product-details	navigation	80.0%	0.0%	0.0%	0.0%
02-rn-nav-stack-filter-default	navigation	100.0%	100.0%	100.0%	100.0%
03-rn-nav-stack-header-from-param	navigation	75.0%	0.0%	0.0%	0.0%
04-rn-nav-modal-compose-note	navigation	100.0%	100.0%	100.0%	100.0%
05-rn-nav-tabs-three-sections	navigation	100.0%	100.0%	100.0%	100.0%
06-rn-nav-home-tab-nested-details	navigation	73.3%	30.0%	91.7%	100.0%
07-rn-nav-drawer-account-help	navigation	100.0%	100.0%	100.0%	100.0%
09-rn-nav-login-replace-home	navigation	100.0%	100.0%	100.0%	100.0%
10-rn-nav-deeplink-profile-basic	navigation	100.0%	100.0%	100.0%	100.0%
12-rn-nav-unsaved-edit-confirm	navigation	100.0%	100.0%	100.0%	100.0%
13-rn-nav-focus-refresh-inbox	navigation	100.0%	100.0%	100.0%	100.0%
14-rn-nav-tabs-lazy-first-render	navigation	100.0%	100.0%	100.0%	100.0%
15-rn-nav-header-action-stateful	navigation	100.0%	100.0%	100.0%	100.0%
01-rn-stylesheet-card-variants	react native apis	100.0%	100.0%	100.0%	100.0%
02-rn-stylesheet-absolute-fill-overlay-badge	react native apis	36.7%	0.0%	0.0%	0.0%
03-rn-stylesheet-divider-hairline	react native apis	100.0%	100.0%	100.0%	100.0%
04-rn-image-network-dimensions	react native apis	100.0%	100.0%	100.0%	100.0%
05-rn-imagebackground-resize-mode	react native apis	100.0%	100.0%	100.0%	100.0%
06-rn-platform-select-small-style-delta	react native apis	100.0%	100.0%	100.0%	100.0%
07-rn-platform-os-conditional-cta	react native apis	100.0%	100.0%	100.0%	100.0%
08-rn-platform-file-extension-component-split	react native apis	100.0%	100.0%	100.0%	100.0%
09-rn-keyboardavoidingview-login-form	react native apis	100.0%	100.0%	100.0%	100.0%

Table 7: Per-eval results (score and pass@k) for claude-opus-4.6

Benchmark (10 runs) · vercel/anthropic/claude-opus-4.6 · score 84.1% (2296/2730 reqs) · pass@1/5/10: 60.6% / 66.8% / 69.7% · cost \$26.07



Figure 3: Benchmark Plot for claude-opus-4.6.

vercel/anthropic/claude-opus-4.6 · Run → Category → Pass/Fail




Figure 4: Sankey per run for claude-opus-4.6 showing the distribution of evals with respect to categories and results

A.2 claude-opus-4.7

Model	Runs	Score	Reqs	P@1	P@5	P@10
vercel/anthropic/claude-opus-4.7	10	82.8%	2260/2730	59.2%	68.3%	71.2%

Table 8: Model summary for claude-opus-4.7

Category	Score	Evals	Requirements	Failed requirements
animation	73.3%	130	53	152
async state	78.9%	130	65	135
lists	83.7%	180	79	137
navigation	95.3%	130	45	24
react native apis	91.9%	90	31	22

Table 9: Results by category for claude-opus-4.7. Score: mean(per-eval passed/total weight) across evals and runs. Evals: count of evaluation tasks. Requirements: sum of requirement count per eval. Failed requirements: sum of failed requirement instances across runs.

Eval	Category	Score	P@1	P@5	P@10
01-rn-anim-pressable-scale-with-timing	animation	100.0%	100.0%	100.0%	100.0%
02-rn-anim-spring-toggle-card-height	animation	90.0%	70.0%	100.0%	100.0%
03-rn-anim-layout-enter-exit-list-items	animation	100.0%	100.0%	100.0%	100.0%
04-rn-anim-scroll-linked-collapsing-header	animation	100.0%	100.0%	100.0%	100.0%
05-rn-anim-pan-drag-card-with-snap-points	animation	75.0%	0.0%	0.0%	0.0%
06-rn-anim-longpress-then-pan-activation-gate	animation	67.5%	30.0%	91.7%	100.0%
07-rn-anim-pinch-pan-simultaneous-photo-canvas	animation	100.0%	100.0%	100.0%	100.0%
09-rn-anim-android-modal-gesture-root-safety	animation	51.7%	0.0%	0.0%	0.0%
10-rn-anim-worklet-thread-bridge-schedule-on-rn	animation	50.0%	0.0%	0.0%	0.0%
11-rn-rngh-exclusive-single-double-tap	animation	58.0%	0.0%	0.0%	0.0%
12-rn-rngh-scroll-pan-failure-chain	animation	76.0%	10.0%	50.0%	100.0%
13-rn-worklets-runtime-offload-and-bridge	animation	33.3%	0.0%	0.0%	0.0%
14-rn-keyboard-controller-aware-scroll-footer	animation	52.0%	0.0%	0.0%	0.0%
01-rn-rq-stable-query-key-filters	async state	40.0%	0.0%	0.0%	0.0%
02-rn-rq-dependent-query-enabled-gate	async state	20.0%	0.0%	0.0%	0.0%
03-rn-rq-mutation-invalidate-on-success	async state	80.0%	0.0%	0.0%	0.0%
04-rn-rq-optimistic-update-rollback	async state	94.3%	70.0%	100.0%	100.0%
05-rn-zustand-async-action-lifecycle	async state	100.0%	100.0%	100.0%	100.0%
06-rn-zustand-request-id-stale-response-guard	async state	80.0%	0.0%	0.0%	0.0%
07-rn-zustand-persist-async-hydration-gate	async state	70.0%	10.0%	50.0%	100.0%
09-rn-jotai-async-read-atom-suspense	async state	100.0%	100.0%	100.0%	100.0%
10-rn-jotai-loadable-async-atom-no-suspense	async state	75.0%	0.0%	0.0%	0.0%
11-rn-jotai-async-write-atom-mutation	async state	100.0%	100.0%	100.0%	100.0%
12-rn-jotai-atomwithstorage-async-rehydrate-guard	async state	100.0%	100.0%	100.0%	100.0%
13-rn-react-usedeferredvalue-search-stale-indicator	async state	100.0%	100.0%	100.0%	100.0%
14-rn-rq-react-native-appstate-focus-online-managers	async state	66.7%	0.0%	0.0%	0.0%
01-rn-flatlist-scroll-to-unread-fixed-rows	lists	100.0%	100.0%	100.0%	100.0%
02-rn-flatlist-basic-contact-list	lists	100.0%	100.0%	100.0%	100.0%
03-rn-flatlist-two-column-product-grid	lists	100.0%	100.0%	100.0%	100.0%
04-rn-flatlist-viewability-impression-tracking	lists	80.0%	0.0%	0.0%	0.0%
05-rn-flatlist-large-feed-window-tuning	lists	100.0%	100.0%	100.0%	100.0%
06-rn-flatlist-pull-to-refresh-feed	lists	90.0%	50.0%	99.6%	100.0%
07-rn-flatlist-to-sectionlist-conversion	lists	100.0%	100.0%	100.0%	100.0%
08-rn-flashlist-basic-list	lists	100.0%	100.0%	100.0%	100.0%
09-rn-flashlist-horizontal-rail-list-header	lists	87.5%	50.0%	99.6%	100.0%
10-rn-flashlist-horizontal-dynamic-height-header-first-item	lists	65.0%	30.0%	91.7%	100.0%
11-rn-flashlist-expandable-cards-layout-state	lists	66.7%	0.0%	0.0%	0.0%
12-rn-flashlist-chat-start-from-bottom	lists	40.0%	0.0%	0.0%	0.0%
13-rn-flashlist-masonry-photo-gallery	lists	40.0%	0.0%	0.0%	0.0%
14-rn-flatlist-to-legendlist-migration	lists	90.0%	70.0%	100.0%	100.0%
15-rn-legendlist-sticky-grouped-agenda	lists	82.0%	10.0%	50.0%	100.0%
16-rn-legendlist-section-scroll-to-location	lists	100.0%	100.0%	100.0%	100.0%
17-rn-legendlist-chat-start-from-bottom	lists	100.0%	100.0%	100.0%	100.0%
18-rn-legendlist-expandable-cards-recycling-state	lists	65.0%	0.0%	0.0%	0.0%
01-rn-nav-stack-product-details	navigation	80.0%	0.0%	0.0%	0.0%
02-rn-nav-stack-filter-default	navigation	100.0%	100.0%	100.0%	100.0%
03-rn-nav-stack-header-from-param	navigation	100.0%	100.0%	100.0%	100.0%
04-rn-nav-modal-compose-note	navigation	100.0%	100.0%	100.0%	100.0%
05-rn-nav-tabs-three-sections	navigation	100.0%	100.0%	100.0%	100.0%
06-rn-nav-home-tab-nested-details	navigation	83.3%	50.0%	99.6%	100.0%
07-rn-nav-drawer-account-help	navigation	100.0%	100.0%	100.0%	100.0%
09-rn-nav-login-replace-home	navigation	100.0%	100.0%	100.0%	100.0%
10-rn-nav-deeplink-profile-basic	navigation	100.0%	100.0%	100.0%	100.0%
12-rn-nav-unsaved-edit-confirm	navigation	80.0%	20.0%	77.8%	100.0%
13-rn-nav-focus-refresh-inbox	navigation	100.0%	100.0%	100.0%	100.0%
14-rn-nav-tabs-lazy-first-render	navigation	100.0%	100.0%	100.0%	100.0%
15-rn-nav-header-action-stateful	navigation	95.0%	90.0%	100.0%	100.0%
01-rn-stylesheet-card-variants	react native apis	66.7%	50.0%	99.6%	100.0%
02-rn-stylesheet-absolute-fill-overlay-badge	react native apis	60.0%	0.0%	0.0%	0.0%
03-rn-stylesheet-divider-hairline	react native apis	100.0%	100.0%	100.0%	100.0%
04-rn-image-network-dimensions	react native apis	100.0%	100.0%	100.0%	100.0%
05-rn-imagebackground-resize-mode	react native apis	100.0%	100.0%	100.0%	100.0%
06-rn-platform-select-small-style-delta	react native apis	100.0%	100.0%	100.0%	100.0%
07-rn-platform-os-conditional-cta	react native apis	100.0%	100.0%	100.0%	100.0%
08-rn-platform-file-extension-component-split	react native apis	100.0%	100.0%	100.0%	100.0%
09-rn-keyboardavoidingview-login-form	react native apis	100.0%	100.0%	100.0%	100.0%

Table 10: Per-eval results (score and pass@k) for claude-opus-4.7

Benchmark (10 runs) · vercel/anthropic/claude-opus-4.7 · score 82.8% (2260/2730 reqs) · pass@1/5/10: 59.2% / 68.3% / 71.2% · cost \$32.27




Figure 5: Benchmark Plot for claude-opus-4.7.

vercel/anthropic/claude-opus-4.7 · Run → Category → Pass/Fail




Figure 6: Sankey per run for claude-opus-4.7 showing the distribution of evals with respect to categories and results

A.3 claude-sonnet-4.6

Model	Runs	Score	Reqs	P@1	P@5	P@10
vercel/anthropic/claude-sonnet-4.6	10	80.6%	2201/2730	55.0%	63.4%	66.7%

Table 11: Model summary for claude-sonnet-4.6

Category	Score	Evals	Requirements	Failed requirements
animation	66.9%	130	53	191
async state	76.9%	130	65	139
lists	80.4%	180	79	151
navigation	93.8%	130	45	30
react native apis	93.3%	90	31	18

Table 12: Results by category for claude-sonnet-4.6. Score: mean(per-eval passed/total weight) across evals and runs. Evals: count of evaluation tasks. Requirements: sum of requirement count per eval. Failed requirements: sum of failed requirement instances across runs.

Eval	Category	Score	P@1	P@5	P@10
01-rn-anim-pressable-scale-with-timing	animation	100.0%	100.0%	100.0%	100.0%
02-rn-anim-spring-toggle-card-height	animation	80.0%	50.0%	99.6%	100.0%
03-rn-anim-layout-enter-exit-list-items	animation	100.0%	100.0%	100.0%	100.0%
04-rn-anim-scroll-linked-collapsing-header	animation	86.7%	60.0%	100.0%	100.0%
05-rn-anim-pan-drag-card-with-snap-points	animation	70.0%	0.0%	0.0%	0.0%
06-rn-anim-longpress-then-pan-activation-gate	animation	87.5%	50.0%	99.6%	100.0%
07-rn-anim-pinch-pan-simultaneous-photo-canvas	animation	100.0%	100.0%	100.0%	100.0%
09-rn-anim-android-modal-gesture-root-safety	animation	40.0%	0.0%	0.0%	0.0%
10-rn-anim-worklet-thread-bridge-schedule-on-rn	animation	40.0%	0.0%	0.0%	0.0%
11-rn-rngh-exclusive-single-double-tap	animation	54.0%	0.0%	0.0%	0.0%
12-rn-rngh-scroll-pan-failure-chain	animation	52.0%	0.0%	0.0%	0.0%
13-rn-worklets-runtime-offload-and-bridge	animation	33.3%	0.0%	0.0%	0.0%
14-rn-keyboard-controller-aware-scroll-footer	animation	26.0%	0.0%	0.0%	0.0%
01-rn-rq-stable-query-key-filters	async state	46.0%	0.0%	0.0%	0.0%
02-rn-rq-dependent-query-enabled-gate	async state	36.0%	0.0%	0.0%	0.0%
03-rn-rq-mutation-invalidate-on-success	async state	72.0%	0.0%	0.0%	0.0%
04-rn-rq-optimistic-update-rollback	async state	95.7%	90.0%	100.0%	100.0%
05-rn-zustand-async-action-lifecycle	async state	100.0%	100.0%	100.0%	100.0%
06-rn-zustand-request-id-stale-response-guard	async state	80.0%	0.0%	0.0%	0.0%
07-rn-zustand-persist-async-hydration-gate	async state	80.0%	10.0%	50.0%	100.0%
09-rn-jotai-async-read-atom-suspense	async state	90.0%	80.0%	100.0%	100.0%
10-rn-jotai-loadable-async-atom-no-suspense	async state	65.0%	0.0%	0.0%	0.0%
11-rn-jotai-async-write-atom-mutation	async state	100.0%	100.0%	100.0%	100.0%
12-rn-jotai-atomwithstorage-async-rehydrate-guard	async state	95.0%	70.0%	100.0%	100.0%
13-rn-react-usedeferredvalue-search-stale-indicator	async state	100.0%	100.0%	100.0%	100.0%
14-rn-rq-react-native-appstate-focus-online-managers	async state	40.0%	0.0%	0.0%	0.0%
01-rn-flatlist-scroll-to-unread-fixed-rows	lists	92.0%	80.0%	100.0%	100.0%
02-rn-flatlist-basic-contact-list	lists	100.0%	100.0%	100.0%	100.0%
03-rn-flatlist-two-column-product-grid	lists	100.0%	100.0%	100.0%	100.0%
04-rn-flatlist-viewability-impression-tracking	lists	80.0%	0.0%	0.0%	0.0%
05-rn-flatlist-large-feed-window-tuning	lists	100.0%	100.0%	100.0%	100.0%
06-rn-flatlist-pull-to-refresh-feed	lists	100.0%	100.0%	100.0%	100.0%
07-rn-flatlist-to-sectionlist-conversion	lists	100.0%	100.0%	100.0%	100.0%
08-rn-flashlist-basic-list	lists	100.0%	100.0%	100.0%	100.0%
09-rn-flashlist-horizontal-rail-list-header	lists	82.5%	30.0%	91.7%	100.0%
10-rn-flashlist-horizontal-dynamic-height-header-first-item	lists	88.3%	30.0%	91.7%	100.0%
11-rn-flashlist-expandable-cards-layout-state	lists	53.3%	0.0%	0.0%	0.0%
12-rn-flashlist-chat-start-from-bottom	lists	40.0%	0.0%	0.0%	0.0%
13-rn-flashlist-masonry-photo-gallery	lists	54.0%	0.0%	0.0%	0.0%
14-rn-flatlist-to-legendlist-migration	lists	46.7%	0.0%	0.0%	0.0%
15-rn-legendlist-sticky-grouped-agenda	lists	80.0%	0.0%	0.0%	0.0%
16-rn-legendlist-section-scroll-to-location	lists	100.0%	100.0%	100.0%	100.0%
17-rn-legendlist-chat-start-from-bottom	lists	80.0%	0.0%	0.0%	0.0%
18-rn-legendlist-expandable-cards-recycling-state	lists	50.0%	0.0%	0.0%	0.0%
01-rn-nav-stack-product-details	navigation	82.0%	10.0%	50.0%	100.0%
02-rn-nav-stack-filter-default	navigation	100.0%	100.0%	100.0%	100.0%
03-rn-nav-stack-header-from-param	navigation	77.5%	10.0%	50.0%	100.0%
04-rn-nav-modal-compose-note	navigation	100.0%	100.0%	100.0%	100.0%
05-rn-nav-tabs-three-sections	navigation	100.0%	100.0%	100.0%	100.0%
06-rn-nav-home-tab-nested-details	navigation	70.0%	10.0%	50.0%	100.0%
07-rn-nav-drawer-account-help	navigation	100.0%	100.0%	100.0%	100.0%
09-rn-nav-login-replace-home	navigation	100.0%	100.0%	100.0%	100.0%
10-rn-nav-deeplink-profile-basic	navigation	100.0%	100.0%	100.0%	100.0%
12-rn-nav-unsaved-edit-confirm	navigation	95.0%	80.0%	100.0%	100.0%
13-rn-nav-focus-refresh-inbox	navigation	100.0%	100.0%	100.0%	100.0%
14-rn-nav-tabs-lazy-first-render	navigation	100.0%	100.0%	100.0%	100.0%
15-rn-nav-header-action-stateful	navigation	95.0%	90.0%	100.0%	100.0%
01-rn-stylesheet-card-variants	react native apis	100.0%	100.0%	100.0%	100.0%
02-rn-stylesheet-absolute-fill-overlay-badge	react native apis	50.0%	0.0%	0.0%	0.0%
03-rn-stylesheet-divider-hairline	react native apis	90.0%	80.0%	100.0%	100.0%
04-rn-image-network-dimensions	react native apis	100.0%	100.0%	100.0%	100.0%
05-rn-imagebackground-resize-mode	react native apis	100.0%	100.0%	100.0%	100.0%
06-rn-platform-select-small-style-delta	react native apis	100.0%	100.0%	100.0%	100.0%
07-rn-platform-os-conditional-cta	react native apis	100.0%	100.0%	100.0%	100.0%
08-rn-platform-file-extension-component-split	react native apis	100.0%	100.0%	100.0%	100.0%
09-rn-keyboardavoidingview-login-form	react native apis	100.0%	100.0%	100.0%	100.0%

Table 13: Per-eval results (score and pass@k) for claude-sonnet-4.6



Figure 7: Benchmark Plot for claude-sonnet-4.6.

vercel/anthropic/claude-sonnet-4.6 · Run → Category → Pass/Fail




Figure 8: Sankey per run for claude-sonnet-4.6 showing the distribution of evals with respect to categories and results

A.4 composer-2

Model	Runs	Score	Reqs	P@1	P@5	P@10
composer-2	10	96.1%	2623/2730	88.2%	96.5%	97.0%

Table 14: Model summary for composer-2

Category	Score	Evals	Requirements	Failed requirements
animation	92.2%	130	53	37
async state	97.3%	130	65	13
lists	96.6%	180	79	23
navigation	97.4%	130	45	12
react native apis	91.9%	90	31	22

Table 15: Results by category for composer-2. Score: mean(per-eval passed/total weight) across evals and runs. Evals: count of evaluation tasks. Requirements: sum of requirement count per eval. Failed requirements: sum of failed requirement instances across runs.

Eval	Category	Score	P@1	P@5	P@10
01-rn-anim-pressable-scale-with-timing	animation	100.0%	100.0%	100.0%	100.0%
02-rn-anim-spring-toggle-card-height	animation	73.3%	30.0%	91.7%	100.0%
03-rn-anim-layout-enter-exit-list-items	animation	100.0%	100.0%	100.0%	100.0%
04-rn-anim-scroll-linked-collapsing-header	animation	80.0%	60.0%	100.0%	100.0%
05-rn-anim-pan-drag-card-with-snap-points	animation	80.0%	20.0%	77.8%	100.0%
06-rn-anim-longpress-then-pan-activation-gate	animation	75.0%	50.0%	99.6%	100.0%
07-rn-anim-pinch-pan-simultaneous-photo-canvas	animation	100.0%	100.0%	100.0%	100.0%
09-rn-anim-android-modal-gesture-root-safety	animation	98.3%	90.0%	100.0%	100.0%
10-rn-anim-worklet-thread-bridge-schedule-on-rn	animation	100.0%	100.0%	100.0%	100.0%
11-rn-rngh-exclusive-single-double-tap	animation	98.0%	90.0%	100.0%	100.0%
12-rn-rngh-scroll-pan-failure-chain	animation	94.0%	90.0%	100.0%	100.0%
13-rn-worklets-runtime-offload-and-bridge	animation	100.0%	100.0%	100.0%	100.0%
14-rn-keyboard-controller-aware-scroll-footer	animation	100.0%	100.0%	100.0%	100.0%
01-rn-rq-stable-query-key-filters	async state	94.0%	90.0%	100.0%	100.0%
02-rn-rq-dependent-query-enabled-gate	async state	100.0%	100.0%	100.0%	100.0%
03-rn-rq-mutation-invalidate-on-success	async state	100.0%	100.0%	100.0%	100.0%
04-rn-rq-optimistic-update-rollback	async state	100.0%	100.0%	100.0%	100.0%
05-rn-zustand-async-action-lifecycle	async state	100.0%	100.0%	100.0%	100.0%
06-rn-zustand-request-id-stale-response-guard	async state	98.0%	90.0%	100.0%	100.0%
07-rn-zustand-persist-async-hydration-gate	async state	100.0%	100.0%	100.0%	100.0%
09-rn-jotai-async-read-atom-suspense	async state	100.0%	100.0%	100.0%	100.0%
10-rn-jotai-loadable-async-atom-no-suspense	async state	97.5%	90.0%	100.0%	100.0%
11-rn-jotai-async-write-atom-mutation	async state	100.0%	100.0%	100.0%	100.0%
12-rn-jotai-atomwithstorage-async-rehydrate-guard	async state	98.3%	90.0%	100.0%	100.0%
13-rn-react-usedeferredvalue-search-stale-indicator	async state	100.0%	100.0%	100.0%	100.0%
14-rn-rq-react-native-appstate-focus-online-managers	async state	76.7%	40.0%	97.6%	100.0%
01-rn-flatlist-scroll-to-unread-fixed-rows	lists	100.0%	100.0%	100.0%	100.0%
02-rn-flatlist-basic-contact-list	lists	100.0%	100.0%	100.0%	100.0%
03-rn-flatlist-two-column-product-grid	lists	100.0%	100.0%	100.0%	100.0%
04-rn-flatlist-viewability-impression-tracking	lists	90.0%	50.0%	99.6%	100.0%
05-rn-flatlist-large-feed-window-tuning	lists	100.0%	100.0%	100.0%	100.0%
06-rn-flatlist-pull-to-refresh-feed	lists	98.0%	90.0%	100.0%	100.0%
07-rn-flatlist-to-sectionlist-conversion	lists	100.0%	100.0%	100.0%	100.0%
08-rn-flashlist-basic-list	lists	100.0%	100.0%	100.0%	100.0%
09-rn-flashlist-horizontal-rail-list-header	lists	100.0%	100.0%	100.0%	100.0%
10-rn-flashlist-horizontal-dynamic-height-header-first-item	lists	100.0%	100.0%	100.0%	100.0%
11-rn-flashlist-expandable-cards-layout-state	lists	66.7%	0.0%	0.0%	0.0%
12-rn-flashlist-chat-start-from-bottom	lists	94.0%	90.0%	100.0%	100.0%
13-rn-flashlist-masonry-photo-gallery	lists	100.0%	100.0%	100.0%	100.0%
14-rn-flatlist-to-legendlist-migration	lists	100.0%	100.0%	100.0%	100.0%
15-rn-legendlist-sticky-grouped-agenda	lists	100.0%	100.0%	100.0%	100.0%
16-rn-legendlist-section-scroll-to-location	lists	100.0%	100.0%	100.0%	100.0%
17-rn-legendlist-chat-start-from-bottom	lists	100.0%	100.0%	100.0%	100.0%
18-rn-legendlist-expandable-cards-recycling-state	lists	90.0%	80.0%	100.0%	100.0%
01-rn-nav-stack-product-details	navigation	100.0%	100.0%	100.0%	100.0%
02-rn-nav-stack-filter-default	navigation	100.0%	100.0%	100.0%	100.0%
03-rn-nav-stack-header-from-param	navigation	100.0%	100.0%	100.0%	100.0%
04-rn-nav-modal-compose-note	navigation	100.0%	100.0%	100.0%	100.0%
05-rn-nav-tabs-three-sections	navigation	100.0%	100.0%	100.0%	100.0%
06-rn-nav-home-tab-nested-details	navigation	80.0%	50.0%	99.6%	100.0%
07-rn-nav-drawer-account-help	navigation	100.0%	100.0%	100.0%	100.0%
09-rn-nav-login-replace-home	navigation	94.0%	80.0%	100.0%	100.0%
10-rn-nav-deeplink-profile-basic	navigation	100.0%	100.0%	100.0%	100.0%
12-rn-nav-unsaved-edit-confirm	navigation	92.5%	70.0%	100.0%	100.0%
13-rn-nav-focus-refresh-inbox	navigation	100.0%	100.0%	100.0%	100.0%
14-rn-nav-tabs-lazy-first-render	navigation	100.0%	100.0%	100.0%	100.0%
15-rn-nav-header-action-stateful	navigation	100.0%	100.0%	100.0%	100.0%
01-rn-stylesheet-card-variants	react native apis	100.0%	100.0%	100.0%	100.0%
02-rn-stylesheet-absolute-fill-overlay-badge	react native apis	33.3%	0.0%	0.0%	0.0%
03-rn-stylesheet-divider-hairline	react native apis	93.3%	80.0%	100.0%	100.0%
04-rn-image-network-dimensions	react native apis	100.0%	100.0%	100.0%	100.0%
05-rn-imagebackground-resize-mode	react native apis	100.0%	100.0%	100.0%	100.0%
06-rn-platform-select-small-style-delta	react native apis	100.0%	100.0%	100.0%	100.0%
07-rn-platform-os-conditional-cta	react native apis	100.0%	100.0%	100.0%	100.0%
08-rn-platform-file-extension-component-split	react native apis	100.0%	100.0%	100.0%	100.0%
09-rn-keyboardavoidingview-login-form	react native apis	100.0%	100.0%	100.0%	100.0%

Table 16: Per-eval results (score and pass@k) for composer-2




Figure 9: Benchmark Plot for composer-2.

composer-2 · Run → Category → Pass/Fail




Figure 10: Sankey per run for composer-2 showing the distribution of evals with respect to categories and results

A.5 composer-2-fast

Model	Runs	Score	Reqs	P@1	P@5	P@10
cursor-default	10	94.9%	2590/2730	85.2%	95.7%	97.0%

Table 17: Model summary for composer-2-fast

Category	Score	Evals	Requirements	Failed requirements
animation	91.8%	130	53	41
async state	96.7%	130	65	16
lists	95.6%	180	79	32
navigation	95.2%	130	45	22
react native apis	89.7%	90	31	29

Table 18: Results by category for composer-2-fast. Score: mean(per-eval passed/total weight) across evals and runs. Evals: count of evaluation tasks. Requirements: sum of requirement count per eval. Failed requirements: sum of failed requirement instances across runs.

Eval	Category	Score	P@1	P@5	P@10
01-rn-anim-pressable-scale-with-timing	animation	100.0%	100.0%	100.0%	100.0%
02-rn-anim-spring-toggle-card-height	animation	66.7%	30.0%	91.7%	100.0%
03-rn-anim-layout-enter-exit-list-items	animation	100.0%	100.0%	100.0%	100.0%
04-rn-anim-scroll-linked-collapsing-header	animation	93.3%	90.0%	100.0%	100.0%
05-rn-anim-pan-drag-card-with-snap-points	animation	82.5%	40.0%	97.6%	100.0%
06-rn-anim-longpress-then-pan-activation-gate	animation	70.0%	40.0%	97.6%	100.0%
07-rn-anim-pinch-pan-simultaneous-photo-canvas	animation	100.0%	100.0%	100.0%	100.0%
09-rn-anim-android-modal-gesture-root-safety	animation	98.3%	90.0%	100.0%	100.0%
10-rn-anim-worklet-thread-bridge-schedule-on-rn	animation	100.0%	100.0%	100.0%	100.0%
11-rn-rngh-exclusive-single-double-tap	animation	86.0%	60.0%	100.0%	100.0%
12-rn-rngh-scroll-pan-failure-chain	animation	96.0%	90.0%	100.0%	100.0%
13-rn-worklets-runtime-offload-and-bridge	animation	100.0%	100.0%	100.0%	100.0%
14-rn-keyboard-controller-aware-scroll-footer	animation	100.0%	100.0%	100.0%	100.0%
01-rn-rq-stable-query-key-filters	async state	94.0%	90.0%	100.0%	100.0%
02-rn-rq-dependent-query-enabled-gate	async state	100.0%	100.0%	100.0%	100.0%
03-rn-rq-mutation-invalidate-on-success	async state	100.0%	100.0%	100.0%	100.0%
04-rn-rq-optimistic-update-rollback	async state	100.0%	100.0%	100.0%	100.0%
05-rn-zustand-async-action-lifecycle	async state	100.0%	100.0%	100.0%	100.0%
06-rn-zustand-request-id-stale-response-guard	async state	96.0%	90.0%	100.0%	100.0%
07-rn-zustand-persist-async-hydration-gate	async state	100.0%	100.0%	100.0%	100.0%
09-rn-jotai-async-read-atom-suspense	async state	100.0%	100.0%	100.0%	100.0%
10-rn-jotai-loadable-async-atom-no-suspense	async state	95.0%	80.0%	100.0%	100.0%
11-rn-jotai-async-write-atom-mutation	async state	100.0%	100.0%	100.0%	100.0%
12-rn-jotai-atomwithstorage-async-rehydrate-guard	async state	98.3%	90.0%	100.0%	100.0%
13-rn-react-usedeferredvalue-search-stale-indicator	async state	97.5%	90.0%	100.0%	100.0%
14-rn-rq-react-native-appstate-focus-online-managers	async state	76.7%	50.0%	99.6%	100.0%
01-rn-flatlist-scroll-to-unread-fixed-rows	lists	96.0%	90.0%	100.0%	100.0%
02-rn-flatlist-basic-contact-list	lists	100.0%	100.0%	100.0%	100.0%
03-rn-flatlist-two-column-product-grid	lists	100.0%	100.0%	100.0%	100.0%
04-rn-flatlist-viewability-impression-tracking	lists	84.0%	20.0%	77.8%	100.0%
05-rn-flatlist-large-feed-window-tuning	lists	98.0%	90.0%	100.0%	100.0%
06-rn-flatlist-pull-to-refresh-feed	lists	98.0%	90.0%	100.0%	100.0%
07-rn-flatlist-to-sectionlist-conversion	lists	100.0%	100.0%	100.0%	100.0%
08-rn-flashlist-basic-list	lists	100.0%	100.0%	100.0%	100.0%
09-rn-flashlist-horizontal-rail-list-header	lists	100.0%	100.0%	100.0%	100.0%
10-rn-flashlist-horizontal-dynamic-height-header-first-item	lists	100.0%	100.0%	100.0%	100.0%
11-rn-flashlist-expandable-cards-layout-state	lists	70.0%	10.0%	50.0%	100.0%
12-rn-flashlist-chat-start-from-bottom	lists	94.0%	90.0%	100.0%	100.0%
13-rn-flashlist-masonry-photo-gallery	lists	100.0%	100.0%	100.0%	100.0%
14-rn-flatlist-to-legendlist-migration	lists	100.0%	100.0%	100.0%	100.0%
15-rn-legendlist-sticky-grouped-agenda	lists	100.0%	100.0%	100.0%	100.0%
16-rn-legendlist-section-scroll-to-location	lists	100.0%	100.0%	100.0%	100.0%
17-rn-legendlist-chat-start-from-bottom	lists	100.0%	100.0%	100.0%	100.0%
18-rn-legendlist-expandable-cards-recycling-state	lists	80.0%	60.0%	100.0%	100.0%
01-rn-nav-stack-product-details	navigation	100.0%	100.0%	100.0%	100.0%
02-rn-nav-stack-filter-default	navigation	100.0%	100.0%	100.0%	100.0%
03-rn-nav-stack-header-from-param	navigation	100.0%	100.0%	100.0%	100.0%
04-rn-nav-modal-compose-note	navigation	100.0%	100.0%	100.0%	100.0%
05-rn-nav-tabs-three-sections	navigation	100.0%	100.0%	100.0%	100.0%
06-rn-nav-home-tab-nested-details	navigation	60.0%	0.0%	0.0%	0.0%
07-rn-nav-drawer-account-help	navigation	100.0%	100.0%	100.0%	100.0%
09-rn-nav-login-replace-home	navigation	90.0%	70.0%	100.0%	100.0%
10-rn-nav-deeplink-profile-basic	navigation	100.0%	100.0%	100.0%	100.0%
12-rn-nav-unsaved-edit-confirm	navigation	87.5%	50.0%	99.6%	100.0%
13-rn-nav-focus-refresh-inbox	navigation	100.0%	100.0%	100.0%	100.0%
14-rn-nav-tabs-lazy-first-render	navigation	100.0%	100.0%	100.0%	100.0%
15-rn-nav-header-action-stateful	navigation	100.0%	100.0%	100.0%	100.0%
01-rn-stylesheet-card-variants	react native apis	100.0%	100.0%	100.0%	100.0%
02-rn-stylesheet-absolute-fill-overlay-badge	react native apis	33.3%	0.0%	0.0%	0.0%
03-rn-stylesheet-divider-hairline	react native apis	86.7%	70.0%	100.0%	100.0%
04-rn-image-network-dimensions	react native apis	100.0%	100.0%	100.0%	100.0%
05-rn-imagebackground-resize-mode	react native apis	87.5%	50.0%	99.6%	100.0%
06-rn-platform-select-small-style-delta	react native apis	100.0%	100.0%	100.0%	100.0%
07-rn-platform-os-conditional-cta	react native apis	100.0%	100.0%	100.0%	100.0%
08-rn-platform-file-extension-component-split	react native apis	100.0%	100.0%	100.0%	100.0%
09-rn-keyboardavoidingview-login-form	react native apis	100.0%	100.0%	100.0%	100.0%

Table 19: Per-eval results (score and pass@k) for composer-2-fast




Figure 11: Benchmark Plot for composer-2-fast.

cursor-default · Run → Category → Pass/Fail




Figure 12: Sankey per run for composer-2-fast showing the distribution of evals with respect to categories and results

A.6 deepseek-r1-distill-qwen-32b

Model	Runs	Score	Reqs	P@1	P@5	P@10
cloudflare-ai-gateway/workers-ai/@cf/deepseek-ai/d	10	44.4%	1211/2730	19.2%	31.1%	36.4%

Table 20: Model summary for deepseek-r1-distill-qwen-32b

Category	Score	Evals	Requirements	Failed requirements
animation	11.1%	130	53	483
async state	32.7%	130	65	438
lists	55.7%	180	79	356
navigation	57.3%	130	45	190
react native apis	82.1%	90	31	52

Table 21: Results by category for deepseek-r1-distill-qwen-32b. Score: mean(per-eval passed/total weight) across evals and runs. Evals: count of evaluation tasks. Requirements: sum of requirement count per eval. Failed requirements: sum of failed requirement instances across runs.

Eval	Category	Score	P@1	P@5	P@10
01-rn-anim-pressable-scale-with-timing	animation	53.3%	10.0%	50.0%	100.0%
02-rn-anim-spring-toggle-card-height	animation	6.7%	0.0%	0.0%	0.0%
03-rn-anim-layout-enter-exit-list-items	animation	40.0%	10.0%	50.0%	100.0%
04-rn-anim-scroll-linked-collapsing-header	animation	6.7%	0.0%	0.0%	0.0%
05-rn-anim-pan-drag-card-with-snap-points	animation	2.5%	0.0%	0.0%	0.0%
06-rn-anim-longpress-then-pan-activation-gate	animation	7.5%	0.0%	0.0%	0.0%
07-rn-anim-pinch-pan-simultaneous-photo-canvas	animation	4.0%	0.0%	0.0%	0.0%
09-rn-anim-android-modal-gesture-root-safety	animation	5.0%	0.0%	0.0%	0.0%
10-rn-anim-worklet-thread-bridge-schedule-on-rn	animation	0.0%	0.0%	0.0%	0.0%
11-rn-rngh-exclusive-single-double-tap	animation	0.0%	0.0%	0.0%	0.0%
12-rn-rngh-scroll-pan-failure-chain	animation	0.0%	0.0%	0.0%	0.0%
13-rn-worklets-runtime-offload-and-bridge	animation	16.7%	0.0%	0.0%	0.0%
14-rn-keyboard-controller-aware-scroll-footer	animation	2.0%	0.0%	0.0%	0.0%
01-rn-rq-stable-query-key-filters	async state	6.0%	0.0%	0.0%	0.0%
02-rn-rq-dependent-query-enabled-gate	async state	18.0%	0.0%	0.0%	0.0%
03-rn-rq-mutation-invalidate-on-success	async state	2.0%	0.0%	0.0%	0.0%
04-rn-rq-optimistic-update-rollback	async state	11.4%	0.0%	0.0%	0.0%
05-rn-zustand-async-action-lifecycle	async state	86.0%	50.0%	99.6%	100.0%
06-rn-zustand-request-id-stale-response-guard	async state	40.0%	0.0%	0.0%	0.0%
07-rn-zustand-persist-async-hydration-gate	async state	28.6%	0.0%	0.0%	0.0%
09-rn-jotai-async-read-atom-suspense	async state	57.5%	20.0%	77.8%	100.0%
10-rn-jotai-loadable-async-atom-no-suspense	async state	0.0%	0.0%	0.0%	0.0%
11-rn-jotai-async-write-atom-mutation	async state	38.0%	0.0%	0.0%	0.0%
12-rn-jotai-atomwithstorage-async-rehydrate-guard	async state	68.3%	0.0%	0.0%	0.0%
13-rn-react-usedeferredvalue-search-stale-indicator	async state	42.5%	10.0%	50.0%	100.0%
14-rn-rq-react-native-appstate-focus-online-managers	async state	26.7%	0.0%	0.0%	0.0%
01-rn-flatlist-scroll-to-unread-fixed-rows	lists	44.0%	0.0%	0.0%	0.0%
02-rn-flatlist-basic-contact-list	lists	100.0%	100.0%	100.0%	100.0%
03-rn-flatlist-two-column-product-grid	lists	93.3%	80.0%	100.0%	100.0%
04-rn-flatlist-viewability-impression-tracking	lists	66.0%	0.0%	0.0%	0.0%
05-rn-flatlist-large-feed-window-tuning	lists	60.0%	0.0%	0.0%	0.0%
06-rn-flatlist-pull-to-refresh-feed	lists	80.0%	0.0%	0.0%	0.0%
07-rn-flatlist-to-sectionlist-conversion	lists	88.0%	40.0%	97.6%	100.0%
08-rn-flashlist-basic-list	lists	76.7%	50.0%	99.6%	100.0%
09-rn-flashlist-horizontal-rail-list-header	lists	42.5%	0.0%	0.0%	0.0%
10-rn-flashlist-horizontal-dynamic-height-header-first-item	lists	53.3%	0.0%	0.0%	0.0%
11-rn-flashlist-expandable-cards-layout-state	lists	36.7%	0.0%	0.0%	0.0%
12-rn-flashlist-chat-start-from-bottom	lists	34.0%	0.0%	0.0%	0.0%
13-rn-flashlist-masonry-photo-gallery	lists	52.0%	0.0%	0.0%	0.0%
14-rn-flatlist-to-legendlist-migration	lists	20.0%	0.0%	0.0%	0.0%
15-rn-legendlist-sticky-grouped-agenda	lists	48.0%	0.0%	0.0%	0.0%
16-rn-legendlist-section-scroll-to-location	lists	52.0%	10.0%	50.0%	100.0%
17-rn-legendlist-chat-start-from-bottom	lists	28.0%	0.0%	0.0%	0.0%
18-rn-legendlist-expandable-cards-recycling-state	lists	27.5%	0.0%	0.0%	0.0%
01-rn-nav-stack-product-details	navigation	60.0%	0.0%	0.0%	0.0%
02-rn-nav-stack-filter-default	navigation	98.0%	90.0%	100.0%	100.0%
03-rn-nav-stack-header-from-param	navigation	50.0%	0.0%	0.0%	0.0%
04-rn-nav-modal-compose-note	navigation	82.5%	40.0%	97.6%	100.0%
05-rn-nav-tabs-three-sections	navigation	60.0%	30.0%	91.7%	100.0%
06-rn-nav-home-tab-nested-details	navigation	70.0%	30.0%	91.7%	100.0%
07-rn-nav-drawer-account-help	navigation	45.0%	10.0%	50.0%	100.0%
09-rn-nav-login-replace-home	navigation	22.0%	0.0%	0.0%	0.0%
10-rn-nav-deeplink-profile-basic	navigation	25.0%	10.0%	50.0%	100.0%
12-rn-nav-unsaved-edit-confirm	navigation	52.5%	0.0%	0.0%	0.0%
13-rn-nav-focus-refresh-inbox	navigation	100.0%	100.0%	100.0%	100.0%
14-rn-nav-tabs-lazy-first-render	navigation	80.0%	40.0%	97.6%	100.0%
15-rn-nav-header-action-stateful	navigation	0.0%	0.0%	0.0%	0.0%
01-rn-stylesheet-card-variants	react native apis	76.7%	50.0%	99.6%	100.0%
02-rn-stylesheet-absolute-fill-overlay-badge	react native apis	66.7%	0.0%	0.0%	0.0%
03-rn-stylesheet-divider-hairline	react native apis	43.3%	0.0%	0.0%	0.0%
04-rn-image-network-dimensions	react native apis	100.0%	100.0%	100.0%	100.0%
05-rn-imagebackground-resize-mode	react native apis	62.5%	0.0%	0.0%	0.0%
06-rn-platform-select-small-style-delta	react native apis	100.0%	100.0%	100.0%	100.0%
07-rn-platform-os-conditional-cta	react native apis	100.0%	100.0%	100.0%	100.0%
08-rn-platform-file-extension-component-split	react native apis	100.0%	100.0%	100.0%	100.0%
09-rn-keyboardavoidingview-login-form	react native apis	90.0%	90.0%	100.0%	100.0%

Table 22: Per-eval results (score and pass@k) for deepseek-r1-distill-qwen-32b

Benchmark (10 runs) · cloudflare-ai-gateway/workers-ai/@cf/deepseek-ai/deepseek-r1-distill-qwen-32b · score 44.4% (1211/2730 reqs) · pass@1/5/10: 19.2% / 31.1% / 36.4% · cost \$4.66

Eval categories (5)

- animation (13)
- lists (18)
- navigation (13)
- react native apis (9)
- async state (13)




Figure 13: Benchmark Plot for deepseek-r1-distill-qwen-32b.

cloudflare-ai-gateway/workers-ai/@cf/deepseek-ai/deepseek-r1-distill-qwen-32b · Run → Cate




Figure 14: Sankey per run for deepseek-r1-distill-qwen-32b showing the distribution of evals with respect to categories and results

A.7 deepseek-v3.2

Model	Runs	Score	Reqs	P@1	P@5	P@10
vercel/deepseek/deepseek-v3.2	10	71.5%	1945/2721	41.8%	56.3%	59.1%

Table 23: Model summary for deepseek-v3.2

Category	Score	Evals	Requirements	Failed requirements
animation	57.5%	129	53	233
async state	71.0%	130	65	182
lists	68.5%	180	79	245
navigation	80.4%	129	45	88
react native apis	89.9%	90	31	28

Table 24: Results by category for deepseek-v3.2. Score: mean(per-eval passed/total weight) across evals and runs. Evals: count of evaluation tasks. Requirements: sum of requirement count per eval. Failed requirements: sum of failed requirement instances across runs.

Eval	Category	Score	P@1	P@5	P@10
01-rn-anim-pressable-scale-with-timing	animation	70.0%	30.0%	91.7%	100.0%
02-rn-anim-spring-toggle-card-height	animation	66.7%	30.0%	91.7%	100.0%
03-rn-anim-layout-enter-exit-list-items	animation	86.7%	80.0%	100.0%	100.0%
04-rn-anim-scroll-linked-collapsing-header	animation	93.3%	90.0%	100.0%	100.0%
05-rn-anim-pan-drag-card-with-snap-points	animation	52.5%	0.0%	0.0%	0.0%
06-rn-anim-longpress-then-pan-activation-gate	animation	82.5%	50.0%	99.6%	100.0%
07-rn-anim-pinch-pan-simultaneous-photo-canvas	animation	98.0%	90.0%	100.0%	100.0%
09-rn-anim-android-modal-gesture-root-safety	animation	48.3%	0.0%	0.0%	0.0%
10-rn-anim-worklet-thread-bridge-schedule-on-rn	animation	22.2%	0.0%	0.0%	0.0%
11-rn-rngh-exclusive-single-double-tap	animation	30.0%	0.0%	0.0%	0.0%
12-rn-rngh-scroll-pan-failure-chain	animation	60.0%	0.0%	0.0%	0.0%
13-rn-worklets-runtime-offload-and-bridge	animation	20.0%	0.0%	0.0%	0.0%
14-rn-keyboard-controller-aware-scroll-footer	animation	14.0%	0.0%	0.0%	0.0%
01-rn-rq-stable-query-key-filters	async state	34.0%	0.0%	0.0%	0.0%
02-rn-rq-dependent-query-enabled-gate	async state	34.0%	0.0%	0.0%	0.0%
03-rn-rq-mutation-invalidate-on-success	async state	72.0%	0.0%	0.0%	0.0%
04-rn-rq-optimistic-update-rollback	async state	91.4%	60.0%	100.0%	100.0%
05-rn-zustand-async-action-lifecycle	async state	100.0%	100.0%	100.0%	100.0%
06-rn-zustand-request-id-stale-response-guard	async state	78.0%	30.0%	91.7%	100.0%
07-rn-zustand-persist-async-hydration-gate	async state	60.0%	0.0%	0.0%	0.0%
09-rn-jotai-async-read-atom-suspense	async state	92.5%	70.0%	100.0%	100.0%
10-rn-jotai-loadable-async-atom-no-suspense	async state	22.5%	0.0%	0.0%	0.0%
11-rn-jotai-async-write-atom-mutation	async state	94.0%	80.0%	100.0%	100.0%
12-rn-jotai-atomwithstorage-async-rehydrate-guard	async state	91.7%	50.0%	99.6%	100.0%
13-rn-react-usedeferredvalue-search-stale-indicator	async state	92.5%	70.0%	100.0%	100.0%
14-rn-rq-react-native-appstate-focus-online-managers	async state	60.0%	30.0%	91.7%	100.0%
01-rn-flatlist-scroll-to-unread-fixed-rows	lists	44.0%	0.0%	0.0%	0.0%
02-rn-flatlist-basic-contact-list	lists	100.0%	100.0%	100.0%	100.0%
03-rn-flatlist-two-column-product-grid	lists	100.0%	100.0%	100.0%	100.0%
04-rn-flatlist-viewability-impression-tracking	lists	80.0%	0.0%	0.0%	0.0%
05-rn-flatlist-large-feed-window-tuning	lists	84.0%	20.0%	77.8%	100.0%
06-rn-flatlist-pull-to-refresh-feed	lists	80.0%	0.0%	0.0%	0.0%
07-rn-flatlist-to-sectionlist-conversion	lists	88.0%	40.0%	97.6%	100.0%
08-rn-flashlist-basic-list	lists	100.0%	100.0%	100.0%	100.0%
09-rn-flashlist-horizontal-rail-list-header	lists	80.0%	20.0%	77.8%	100.0%
10-rn-flashlist-horizontal-dynamic-height-header-first-item	lists	81.7%	10.0%	50.0%	100.0%
11-rn-flashlist-expandable-cards-layout-state	lists	43.3%	0.0%	0.0%	0.0%
12-rn-flashlist-chat-start-from-bottom	lists	40.0%	0.0%	0.0%	0.0%
13-rn-flashlist-masonry-photo-gallery	lists	66.0%	0.0%	0.0%	0.0%
14-rn-flatlist-to-legendlist-migration	lists	0.0%	0.0%	0.0%	0.0%
15-rn-legendlist-sticky-grouped-agenda	lists	78.0%	0.0%	0.0%	0.0%
16-rn-legendlist-section-scroll-to-location	lists	100.0%	100.0%	100.0%	100.0%
17-rn-legendlist-chat-start-from-bottom	lists	38.0%	0.0%	0.0%	0.0%
18-rn-legendlist-expandable-cards-recycling-state	lists	30.0%	0.0%	0.0%	0.0%
01-rn-nav-stack-product-details	navigation	68.0%	0.0%	0.0%	0.0%
02-rn-nav-stack-filter-default	navigation	100.0%	100.0%	100.0%	100.0%
03-rn-nav-stack-header-from-param	navigation	55.0%	10.0%	50.0%	100.0%
04-rn-nav-modal-compose-note	navigation	100.0%	100.0%	100.0%	100.0%
05-rn-nav-tabs-three-sections	navigation	96.7%	90.0%	100.0%	100.0%
06-rn-nav-home-tab-nested-details	navigation	90.0%	70.0%	100.0%	100.0%
07-rn-nav-drawer-account-help	navigation	100.0%	100.0%	100.0%	100.0%
09-rn-nav-login-replace-home	navigation	57.8%	0.0%	0.0%	0.0%
10-rn-nav-deeplink-profile-basic	navigation	100.0%	100.0%	100.0%	100.0%
12-rn-nav-unsaved-edit-confirm	navigation	75.0%	0.0%	0.0%	0.0%
13-rn-nav-focus-refresh-inbox	navigation	100.0%	100.0%	100.0%	100.0%
14-rn-nav-tabs-lazy-first-render	navigation	90.0%	70.0%	100.0%	100.0%
15-rn-nav-header-action-stateful	navigation	10.0%	0.0%	0.0%	0.0%
01-rn-stylesheet-card-variants	react native apis	66.7%	50.0%	99.6%	100.0%
02-rn-stylesheet-absolute-fill-overlay-badge	react native apis	66.7%	0.0%	0.0%	0.0%
03-rn-stylesheet-divider-hairline	react native apis	83.3%	50.0%	99.6%	100.0%
04-rn-image-network-dimensions	react native apis	100.0%	100.0%	100.0%	100.0%
05-rn-imagebackground-resize-mode	react native apis	92.5%	70.0%	100.0%	100.0%
06-rn-platform-select-small-style-delta	react native apis	100.0%	100.0%	100.0%	100.0%
07-rn-platform-os-conditional-cta	react native apis	100.0%	100.0%	100.0%	100.0%
08-rn-platform-file-extension-component-split	react native apis	100.0%	100.0%	100.0%	100.0%
09-rn-keyboardavoidingview-login-form	react native apis	100.0%	100.0%	100.0%	100.0%

Table 25: Per-eval results (score and pass@k) for deepseek-v3.2




Figure 15: Benchmark Plot for deepseek-v3.2.

vercel/deepseek/deepseek-v3.2 · Run → Category → Pass/Fail




Figure 16: Sankey per run for deepseek-v3.2 showing the distribution of evals with respect to categories and results

A.8 gemini-3.1-pro-preview

Model	Runs	Score	Reqs	P@1	P@5	P@10
vercel/google/gemini-3.1-pro-preview	10	78.9%	2154/2730	52.4%	60.3%	63.6%

Table 26: Model summary for gemini-3.1-pro-preview

Category	Score	Evals	Requirements	Failed requirements
animation	65.1%	130	53	189
async state	78.9%	130	65	130
lists	73.7%	180	79	215
navigation	95.0%	130	45	25
react native apis	93.7%	90	31	17

Table 27: Results by category for gemini-3.1-pro-preview. Score: mean(per-eval passed/total weight) across evals and runs. Evals: count of evaluation tasks. Requirements: sum of requirement count per eval. Failed requirements: sum of failed requirement instances across runs.

Eval	Category	Score	P@1	P@5	P@10
01-rn-anim-pressable-scale-with-timing	animation	100.0%	100.0%	100.0%	100.0%
02-rn-anim-spring-toggle-card-height	animation	40.0%	0.0%	0.0%	0.0%
03-rn-anim-layout-enter-exit-list-items	animation	100.0%	100.0%	100.0%	100.0%
04-rn-anim-scroll-linked-collapsing-header	animation	76.7%	60.0%	100.0%	100.0%
05-rn-anim-pan-drag-card-with-snap-points	animation	75.0%	0.0%	0.0%	0.0%
06-rn-anim-longpress-then-pan-activation-gate	animation	55.0%	10.0%	50.0%	100.0%
07-rn-anim-pinch-pan-simultaneous-photo-canvas	animation	100.0%	100.0%	100.0%	100.0%
09-rn-anim-android-modal-gesture-root-safety	animation	55.0%	0.0%	0.0%	0.0%
10-rn-anim-worklet-thread-bridge-schedule-on-rn	animation	50.0%	0.0%	0.0%	0.0%
11-rn-rngh-exclusive-single-double-tap	animation	58.0%	0.0%	0.0%	0.0%
12-rn-rngh-scroll-pan-failure-chain	animation	78.0%	0.0%	0.0%	0.0%
13-rn-worklets-runtime-offload-and-bridge	animation	30.0%	0.0%	0.0%	0.0%
14-rn-keyboard-controller-aware-scroll-footer	animation	28.0%	0.0%	0.0%	0.0%
01-rn-rq-stable-query-key-filters	async state	40.0%	0.0%	0.0%	0.0%
02-rn-rq-dependent-query-enabled-gate	async state	56.0%	0.0%	0.0%	0.0%
03-rn-rq-mutation-invalidate-on-success	async state	40.0%	0.0%	0.0%	0.0%
04-rn-rq-optimistic-update-rollback	async state	100.0%	100.0%	100.0%	100.0%
05-rn-zustand-async-action-lifecycle	async state	100.0%	100.0%	100.0%	100.0%
06-rn-zustand-request-id-stale-response-guard	async state	78.0%	0.0%	0.0%	0.0%
07-rn-zustand-persist-async-hydration-gate	async state	82.9%	40.0%	97.6%	100.0%
09-rn-jotai-async-read-atom-suspense	async state	100.0%	100.0%	100.0%	100.0%
10-rn-jotai-loadable-async-atom-no-suspense	async state	75.0%	0.0%	0.0%	0.0%
11-rn-jotai-async-write-atom-mutation	async state	100.0%	100.0%	100.0%	100.0%
12-rn-jotai-atomwithstorage-async-rehydrate-guard	async state	98.3%	90.0%	100.0%	100.0%
13-rn-react-usedeferredvalue-search-stale-indicator	async state	92.5%	70.0%	100.0%	100.0%
14-rn-rq-react-native-appstate-focus-online-managers	async state	63.3%	0.0%	0.0%	0.0%
01-rn-flatlist-scroll-to-unread-fixed-rows	lists	100.0%	100.0%	100.0%	100.0%
02-rn-flatlist-basic-contact-list	lists	100.0%	100.0%	100.0%	100.0%
03-rn-flatlist-two-column-product-grid	lists	100.0%	100.0%	100.0%	100.0%
04-rn-flatlist-viewability-impression-tracking	lists	80.0%	0.0%	0.0%	0.0%
05-rn-flatlist-large-feed-window-tuning	lists	84.0%	20.0%	77.8%	100.0%
06-rn-flatlist-pull-to-refresh-feed	lists	84.0%	20.0%	77.8%	100.0%
07-rn-flatlist-to-sectionlist-conversion	lists	100.0%	100.0%	100.0%	100.0%
08-rn-flashlist-basic-list	lists	73.3%	20.0%	77.8%	100.0%
09-rn-flashlist-horizontal-rail-list-header	lists	90.0%	60.0%	100.0%	100.0%
10-rn-flashlist-horizontal-dynamic-height-header-first-item	lists	50.0%	0.0%	0.0%	0.0%
11-rn-flashlist-expandable-cards-layout-state	lists	63.3%	0.0%	0.0%	0.0%
12-rn-flashlist-chat-start-from-bottom	lists	36.0%	0.0%	0.0%	0.0%
13-rn-flashlist-masonry-photo-gallery	lists	30.0%	0.0%	0.0%	0.0%
14-rn-flatlist-to-legendlist-migration	lists	66.7%	0.0%	0.0%	0.0%
15-rn-legendlist-sticky-grouped-agenda	lists	80.0%	0.0%	0.0%	0.0%
16-rn-legendlist-section-scroll-to-location	lists	100.0%	100.0%	100.0%	100.0%
17-rn-legendlist-chat-start-from-bottom	lists	56.0%	0.0%	0.0%	0.0%
18-rn-legendlist-expandable-cards-recycling-state	lists	32.5%	0.0%	0.0%	0.0%
01-rn-nav-stack-product-details	navigation	78.0%	10.0%	50.0%	100.0%
02-rn-nav-stack-filter-default	navigation	100.0%	100.0%	100.0%	100.0%
03-rn-nav-stack-header-from-param	navigation	100.0%	100.0%	100.0%	100.0%
04-rn-nav-modal-compose-note	navigation	100.0%	100.0%	100.0%	100.0%
05-rn-nav-tabs-three-sections	navigation	100.0%	100.0%	100.0%	100.0%
06-rn-nav-home-tab-nested-details	navigation	70.0%	10.0%	50.0%	100.0%
07-rn-nav-drawer-account-help	navigation	100.0%	100.0%	100.0%	100.0%
09-rn-nav-login-replace-home	navigation	100.0%	100.0%	100.0%	100.0%
10-rn-nav-deeplink-profile-basic	navigation	100.0%	100.0%	100.0%	100.0%
12-rn-nav-unsaved-edit-confirm	navigation	87.5%	50.0%	99.6%	100.0%
13-rn-nav-focus-refresh-inbox	navigation	100.0%	100.0%	100.0%	100.0%
14-rn-nav-tabs-lazy-first-render	navigation	100.0%	100.0%	100.0%	100.0%
15-rn-nav-header-action-stateful	navigation	100.0%	100.0%	100.0%	100.0%
01-rn-stylesheet-card-variants	react native apis	100.0%	100.0%	100.0%	100.0%
02-rn-stylesheet-absolute-fill-overlay-badge	react native apis	43.3%	0.0%	0.0%	0.0%
03-rn-stylesheet-divider-hairline	react native apis	100.0%	100.0%	100.0%	100.0%
04-rn-image-network-dimensions	react native apis	100.0%	100.0%	100.0%	100.0%
05-rn-imagebackground-resize-mode	react native apis	100.0%	100.0%	100.0%	100.0%
06-rn-platform-select-small-style-delta	react native apis	100.0%	100.0%	100.0%	100.0%
07-rn-platform-os-conditional-cta	react native apis	100.0%	100.0%	100.0%	100.0%
08-rn-platform-file-extension-component-split	react native apis	100.0%	100.0%	100.0%	100.0%
09-rn-keyboardavoidingview-login-form	react native apis	100.0%	100.0%	100.0%	100.0%

Table 28: Per-eval results (score and pass@k) for gemini-3.1-pro-preview




Figure 17: Benchmark Plot for gemini-3.1-pro-preview.

vercel/google/gemini-3.1-pro-preview · Run → Category → Pass/Fail




Figure 18: Sankey per run for gemini-3.1-pro-preview showing the distribution of evals with respect to categories and results

A.9 gemma-4-31B-it

Model	Runs	Score	Reqs	P@1	P@5	P@10
llama.cpp/google/gemma-4-31B-it	10	75.2%	2054/2730	49.7%	57.7%	59.1%

Table 29: Model summary for gemma-4-31B-it

Category	Score	Evals	Requirements	Failed requirements
animation	65.8%	130	53	198
async state	76.9%	130	65	145
lists	67.8%	180	79	263
navigation	91.1%	130	45	50
react native apis	92.6%	90	31	20

Table 30: Results by category for gemma-4-31B-it. Score: mean(per-eval passed/total weight) across evals and runs. Evals: count of evaluation tasks. Requirements: sum of requirement count per eval. Failed requirements: sum of failed requirement instances across runs.

Eval	Category	Score	P@1	P@5	P@10
01-rn-anim-pressable-scale-with-timing	animation	100.0%	100.0%	100.0%	100.0%
02-rn-anim-spring-toggle-card-height	animation	73.3%	40.0%	97.6%	100.0%
03-rn-anim-layout-enter-exit-list-items	animation	100.0%	100.0%	100.0%	100.0%
04-rn-anim-scroll-linked-collapsing-header	animation	100.0%	100.0%	100.0%	100.0%
05-rn-anim-pan-drag-card-with-snap-points	animation	67.5%	0.0%	0.0%	0.0%
06-rn-anim-longpress-then-pan-activation-gate	animation	80.0%	60.0%	100.0%	100.0%
07-rn-anim-pinch-pan-simultaneous-photo-canvas	animation	100.0%	100.0%	100.0%	100.0%
09-rn-anim-android-modal-gesture-root-safety	animation	45.0%	0.0%	0.0%	0.0%
10-rn-anim-worklet-thread-bridge-schedule-on-rn	animation	50.0%	0.0%	0.0%	0.0%
11-rn-rngh-exclusive-single-double-tap	animation	52.0%	0.0%	0.0%	0.0%
12-rn-rngh-scroll-pan-failure-chain	animation	48.0%	0.0%	0.0%	0.0%
13-rn-worklets-runtime-offload-and-bridge	animation	30.0%	0.0%	0.0%	0.0%
14-rn-keyboard-controller-aware-scroll-footer	animation	10.0%	0.0%	0.0%	0.0%
01-rn-rq-stable-query-key-filters	async state	40.0%	0.0%	0.0%	0.0%
02-rn-rq-dependent-query-enabled-gate	async state	48.0%	0.0%	0.0%	0.0%
03-rn-rq-mutation-invalidate-on-success	async state	72.0%	0.0%	0.0%	0.0%
04-rn-rq-optimistic-update-rollback	async state	100.0%	100.0%	100.0%	100.0%
05-rn-zustand-async-action-lifecycle	async state	100.0%	100.0%	100.0%	100.0%
06-rn-zustand-request-id-stale-response-guard	async state	48.0%	0.0%	0.0%	0.0%
07-rn-zustand-persist-async-hydration-gate	async state	71.4%	0.0%	0.0%	0.0%
09-rn-jotai-async-read-atom-suspense	async state	100.0%	100.0%	100.0%	100.0%
10-rn-jotai-loadable-async-atom-no-suspense	async state	75.0%	0.0%	0.0%	0.0%
11-rn-jotai-async-write-atom-mutation	async state	100.0%	100.0%	100.0%	100.0%
12-rn-jotai-atomwithstorage-async-rehydrate-guard	async state	95.0%	70.0%	100.0%	100.0%
13-rn-react-usedeferredvalue-search-stale-indicator	async state	87.5%	50.0%	99.6%	100.0%
14-rn-rq-react-native-appstate-focus-online-managers	async state	63.3%	0.0%	0.0%	0.0%
01-rn-flatlist-scroll-to-unread-fixed-rows	lists	86.0%	30.0%	91.7%	100.0%
02-rn-flatlist-basic-contact-list	lists	100.0%	100.0%	100.0%	100.0%
03-rn-flatlist-two-column-product-grid	lists	100.0%	100.0%	100.0%	100.0%
04-rn-flatlist-viewability-impression-tracking	lists	80.0%	0.0%	0.0%	0.0%
05-rn-flatlist-large-feed-window-tuning	lists	84.0%	20.0%	77.8%	100.0%
06-rn-flatlist-pull-to-refresh-feed	lists	80.0%	0.0%	0.0%	0.0%
07-rn-flatlist-to-sectionlist-conversion	lists	100.0%	100.0%	100.0%	100.0%
08-rn-flashlist-basic-list	lists	90.0%	70.0%	100.0%	100.0%
09-rn-flashlist-horizontal-rail-list-header	lists	67.5%	30.0%	91.7%	100.0%
10-rn-flashlist-horizontal-dynamic-height-header-first-item	lists	46.7%	0.0%	0.0%	0.0%
11-rn-flashlist-expandable-cards-layout-state	lists	33.3%	0.0%	0.0%	0.0%
12-rn-flashlist-chat-start-from-bottom	lists	20.0%	0.0%	0.0%	0.0%
13-rn-flashlist-masonry-photo-gallery	lists	20.0%	0.0%	0.0%	0.0%
14-rn-flatlist-to-legendlist-migration	lists	70.0%	10.0%	50.0%	100.0%
15-rn-legendlist-sticky-grouped-agenda	lists	76.0%	0.0%	0.0%	0.0%
16-rn-legendlist-section-scroll-to-location	lists	100.0%	100.0%	100.0%	100.0%
17-rn-legendlist-chat-start-from-bottom	lists	40.0%	0.0%	0.0%	0.0%
18-rn-legendlist-expandable-cards-recycling-state	lists	27.5%	0.0%	0.0%	0.0%
01-rn-nav-stack-product-details	navigation	66.0%	0.0%	0.0%	0.0%
02-rn-nav-stack-filter-default	navigation	100.0%	100.0%	100.0%	100.0%
03-rn-nav-stack-header-from-param	navigation	87.5%	50.0%	99.6%	100.0%
04-rn-nav-modal-compose-note	navigation	100.0%	100.0%	100.0%	100.0%
05-rn-nav-tabs-three-sections	navigation	100.0%	100.0%	100.0%	100.0%
06-rn-nav-home-tab-nested-details	navigation	86.7%	60.0%	100.0%	100.0%
07-rn-nav-drawer-account-help	navigation	100.0%	100.0%	100.0%	100.0%
09-rn-nav-login-replace-home	navigation	76.0%	0.0%	0.0%	0.0%
10-rn-nav-deeplink-profile-basic	navigation	100.0%	100.0%	100.0%	100.0%
12-rn-nav-unsaved-edit-confirm	navigation	72.5%	0.0%	0.0%	0.0%
13-rn-nav-focus-refresh-inbox	navigation	100.0%	100.0%	100.0%	100.0%
14-rn-nav-tabs-lazy-first-render	navigation	100.0%	100.0%	100.0%	100.0%
15-rn-nav-header-action-stateful	navigation	95.0%	90.0%	100.0%	100.0%
01-rn-stylesheet-card-variants	react native apis	100.0%	100.0%	100.0%	100.0%
02-rn-stylesheet-absolute-fill-overlay-badge	react native apis	33.3%	0.0%	0.0%	0.0%
03-rn-stylesheet-divider-hairline	react native apis	100.0%	100.0%	100.0%	100.0%
04-rn-image-network-dimensions	react native apis	100.0%	100.0%	100.0%	100.0%
05-rn-imagebackground-resize-mode	react native apis	100.0%	100.0%	100.0%	100.0%
06-rn-platform-select-small-style-delta	react native apis	100.0%	100.0%	100.0%	100.0%
07-rn-platform-os-conditional-cta	react native apis	100.0%	100.0%	100.0%	100.0%
08-rn-platform-file-extension-component-split	react native apis	100.0%	100.0%	100.0%	100.0%
09-rn-keyboardavoidingview-login-form	react native apis	100.0%	100.0%	100.0%	100.0%

Table 31: Per-eval results (score and pass@k) for gemma-4-31B-it




Figure 19: Benchmark Plot for gemma-4-31B-it.

llama.cpp/google/gemma-4-31B-it · Run → Category → Pass/Fail




Figure 20: Sankey per run for gemma-4-31B-it showing the distribution of evals with respect to categories and results

A.10 glm-5

Model	Runs	Score	Reqs	P@1	P@5	P@10
vercel/zai/glm-5	10	74.8%	2043/2730	50.5%	60.3%	60.6%

Table 32: Model summary for glm-5

Category	Score	Evals	Requirements	Failed requirements
animation	66.1%	130	53	194
async state	72.1%	130	65	169
lists	70.7%	180	79	227
navigation	89.2%	130	45	57
react native apis	85.4%	90	31	40

Table 33: Results by category for glm-5. Score: mean(per-eval passed/total weight) across evals and runs. Evals: count of evaluation tasks. Requirements: sum of requirement count per eval. Failed requirements: sum of failed requirement instances across runs.

Eval	Category	Score	P@1	P@5	P@10
01-rn-anim-pressable-scale-with-timing	animation	100.0%	100.0%	100.0%	100.0%
02-rn-anim-spring-toggle-card-height	animation	83.3%	50.0%	99.6%	100.0%
03-rn-anim-layout-enter-exit-list-items	animation	100.0%	100.0%	100.0%	100.0%
04-rn-anim-scroll-linked-collapsing-header	animation	86.7%	70.0%	100.0%	100.0%
05-rn-anim-pan-drag-card-with-snap-points	animation	52.5%	0.0%	0.0%	0.0%
06-rn-anim-longpress-then-pan-activation-gate	animation	82.5%	50.0%	99.6%	100.0%
07-rn-anim-pinch-pan-simultaneous-photo-canvas	animation	98.0%	90.0%	100.0%	100.0%
09-rn-anim-android-modal-gesture-root-safety	animation	46.7%	0.0%	0.0%	0.0%
10-rn-anim-worklet-thread-bridge-schedule-on-rn	animation	50.0%	0.0%	0.0%	0.0%
11-rn-rngh-exclusive-single-double-tap	animation	40.0%	0.0%	0.0%	0.0%
12-rn-rngh-scroll-pan-failure-chain	animation	74.0%	0.0%	0.0%	0.0%
13-rn-worklets-runtime-offload-and-bridge	animation	30.0%	0.0%	0.0%	0.0%
14-rn-keyboard-controller-aware-scroll-footer	animation	16.0%	0.0%	0.0%	0.0%
01-rn-rq-stable-query-key-filters	async state	40.0%	0.0%	0.0%	0.0%
02-rn-rq-dependent-query-enabled-gate	async state	42.0%	0.0%	0.0%	0.0%
03-rn-rq-mutation-invalidate-on-success	async state	58.0%	0.0%	0.0%	0.0%
04-rn-rq-optimistic-update-rollback	async state	98.6%	90.0%	100.0%	100.0%
05-rn-zustand-async-action-lifecycle	async state	100.0%	100.0%	100.0%	100.0%
06-rn-zustand-request-id-stale-response-guard	async state	68.0%	30.0%	91.7%	100.0%
07-rn-zustand-persist-async-hydration-gate	async state	64.3%	0.0%	0.0%	0.0%
09-rn-jotai-async-read-atom-suspense	async state	100.0%	100.0%	100.0%	100.0%
10-rn-jotai-loadable-async-atom-no-suspense	async state	40.0%	0.0%	0.0%	0.0%
11-rn-jotai-async-write-atom-mutation	async state	100.0%	100.0%	100.0%	100.0%
12-rn-jotai-atomwithstorage-async-rehydrate-guard	async state	98.3%	90.0%	100.0%	100.0%
13-rn-react-usedeferredvalue-search-stale-indicator	async state	95.0%	80.0%	100.0%	100.0%
14-rn-rq-react-native-appstate-focus-online-managers	async state	33.3%	0.0%	0.0%	0.0%
01-rn-flatlist-scroll-to-unread-fixed-rows	lists	72.0%	0.0%	0.0%	0.0%
02-rn-flatlist-basic-contact-list	lists	100.0%	100.0%	100.0%	100.0%
03-rn-flatlist-two-column-product-grid	lists	100.0%	100.0%	100.0%	100.0%
04-rn-flatlist-viewability-impression-tracking	lists	80.0%	0.0%	0.0%	0.0%
05-rn-flatlist-large-feed-window-tuning	lists	94.0%	70.0%	100.0%	100.0%
06-rn-flatlist-pull-to-refresh-feed	lists	80.0%	0.0%	0.0%	0.0%
07-rn-flatlist-to-sectionlist-conversion	lists	96.0%	80.0%	100.0%	100.0%
08-rn-flashlist-basic-list	lists	100.0%	100.0%	100.0%	100.0%
09-rn-flashlist-horizontal-rail-list-header	lists	97.5%	90.0%	100.0%	100.0%
10-rn-flashlist-horizontal-dynamic-height-header-first-item	lists	66.7%	0.0%	0.0%	0.0%
11-rn-flashlist-expandable-cards-layout-state	lists	36.7%	0.0%	0.0%	0.0%
12-rn-flashlist-chat-start-from-bottom	lists	40.0%	0.0%	0.0%	0.0%
13-rn-flashlist-masonry-photo-gallery	lists	54.0%	0.0%	0.0%	0.0%
14-rn-flatlist-to-legendlist-migration	lists	0.0%	0.0%	0.0%	0.0%
15-rn-legendlist-sticky-grouped-agenda	lists	88.0%	40.0%	97.6%	100.0%
16-rn-legendlist-section-scroll-to-location	lists	100.0%	100.0%	100.0%	100.0%
17-rn-legendlist-chat-start-from-bottom	lists	40.0%	0.0%	0.0%	0.0%
18-rn-legendlist-expandable-cards-recycling-state	lists	27.5%	0.0%	0.0%	0.0%
01-rn-nav-stack-product-details	navigation	66.0%	0.0%	0.0%	0.0%
02-rn-nav-stack-filter-default	navigation	100.0%	100.0%	100.0%	100.0%
03-rn-nav-stack-header-from-param	navigation	100.0%	100.0%	100.0%	100.0%
04-rn-nav-modal-compose-note	navigation	100.0%	100.0%	100.0%	100.0%
05-rn-nav-tabs-three-sections	navigation	100.0%	100.0%	100.0%	100.0%
06-rn-nav-home-tab-nested-details	navigation	80.0%	40.0%	97.6%	100.0%
07-rn-nav-drawer-account-help	navigation	97.5%	90.0%	100.0%	100.0%
09-rn-nav-login-replace-home	navigation	58.0%	0.0%	0.0%	0.0%
10-rn-nav-deeplink-profile-basic	navigation	100.0%	100.0%	100.0%	100.0%
12-rn-nav-unsaved-edit-confirm	navigation	82.5%	40.0%	97.6%	100.0%
13-rn-nav-focus-refresh-inbox	navigation	100.0%	100.0%	100.0%	100.0%
14-rn-nav-tabs-lazy-first-render	navigation	100.0%	100.0%	100.0%	100.0%
15-rn-nav-header-action-stateful	navigation	75.0%	60.0%	100.0%	100.0%
01-rn-stylesheet-card-variants	react native apis	60.0%	40.0%	97.6%	100.0%
02-rn-stylesheet-absolute-fill-overlay-badge	react native apis	36.7%	0.0%	0.0%	0.0%
03-rn-stylesheet-divider-hairline	react native apis	73.3%	40.0%	97.6%	100.0%
04-rn-image-network-dimensions	react native apis	100.0%	100.0%	100.0%	100.0%
05-rn-imagebackground-resize-mode	react native apis	100.0%	100.0%	100.0%	100.0%
06-rn-platform-select-small-style-delta	react native apis	98.3%	90.0%	100.0%	100.0%
07-rn-platform-os-conditional-cta	react native apis	100.0%	100.0%	100.0%	100.0%
08-rn-platform-file-extension-component-split	react native apis	100.0%	100.0%	100.0%	100.0%
09-rn-keyboardavoidingview-login-form	react native apis	100.0%	100.0%	100.0%	100.0%

Table 34: Per-eval results (score and pass@k) for glm-5




Figure 21: Benchmark Plot for glm-5.

vercel/zai/glm-5 · Run → Category → Pass/Fail




Figure 22: Sankey per run for glm-5 showing the distribution of evals with respect to categories and results

A.11 gpt-5.3-codex

Model	Runs	Score	Reqs	P@1	P@5	P@10
vercel/openai/gpt-5.3-codex	10	83.1%	2267/2727	63.3%	74.5%	78.5%

Table 35: Model summary for gpt-5.3-codex

Category	Score	Evals	Requirements	Failed requirements
animation	67.1%	130	53	186
async state	83.4%	129	65	105
lists	83.8%	180	79	130
navigation	96.2%	130	45	20
react native apis	93.0%	90	31	19

Table 36: Results by category for gpt-5.3-codex. Score: mean(per-eval passed/total weight) across evals and runs. Evals: count of evaluation tasks. Requirements: sum of requirement count per eval. Failed requirements: sum of failed requirement instances across runs.

Eval	Category	Score	P@1	P@5	P@10
01-rn-anim-pressable-scale-with-timing	animation	100.0%	100.0%	100.0%	100.0%
02-rn-anim-spring-toggle-card-height	animation	36.7%	10.0%	50.0%	100.0%
03-rn-anim-layout-enter-exit-list-items	animation	100.0%	100.0%	100.0%	100.0%
04-rn-anim-scroll-linked-collapsing-header	animation	96.7%	90.0%	100.0%	100.0%
05-rn-anim-pan-drag-card-with-snap-points	animation	77.5%	10.0%	50.0%	100.0%
06-rn-anim-longpress-then-pan-activation-gate	animation	72.5%	40.0%	97.6%	100.0%
07-rn-anim-pinch-pan-simultaneous-photo-canvas	animation	100.0%	100.0%	100.0%	100.0%
09-rn-anim-android-modal-gesture-root-safety	animation	23.3%	0.0%	0.0%	0.0%
10-rn-anim-worklet-thread-bridge-schedule-on-rn	animation	47.5%	0.0%	0.0%	0.0%
11-rn-rngh-exclusive-single-double-tap	animation	56.0%	0.0%	0.0%	0.0%
12-rn-rngh-scroll-pan-failure-chain	animation	62.0%	20.0%	77.8%	100.0%
13-rn-worklets-runtime-offload-and-bridge	animation	40.0%	10.0%	50.0%	100.0%
14-rn-keyboard-controller-aware-scroll-footer	animation	60.0%	0.0%	0.0%	0.0%
01-rn-rq-stable-query-key-filters	async state	46.0%	0.0%	0.0%	0.0%
02-rn-rq-dependent-query-enabled-gate	async state	34.0%	0.0%	0.0%	0.0%
03-rn-rq-mutation-invalidate-on-success	async state	58.0%	0.0%	0.0%	0.0%
04-rn-rq-optimistic-update-rollback	async state	98.6%	90.0%	100.0%	100.0%
05-rn-zustand-async-action-lifecycle	async state	100.0%	100.0%	100.0%	100.0%
06-rn-zustand-request-id-stale-response-guard	async state	90.0%	50.0%	99.6%	100.0%
07-rn-zustand-persist-async-hydration-gate	async state	91.4%	60.0%	100.0%	100.0%
09-rn-jotai-async-read-atom-suspense	async state	100.0%	100.0%	100.0%	100.0%
10-rn-jotai-loadable-async-atom-no-suspense	async state	75.0%	0.0%	0.0%	0.0%
11-rn-jotai-async-write-atom-mutation	async state	100.0%	100.0%	100.0%	100.0%
12-rn-jotai-atomwithstorage-async-rehydrate-guard	async state	100.0%	100.0%	100.0%	100.0%
13-rn-react-usedeferredvalue-search-stale-indicator	async state	100.0%	100.0%	100.0%	100.0%
14-rn-rq-react-native-appstate-focus-online-managers	async state	92.6%	77.8%	100.0%	77.8%
01-rn-flatlist-scroll-to-unread-fixed-rows	lists	100.0%	100.0%	100.0%	100.0%
02-rn-flatlist-basic-contact-list	lists	100.0%	100.0%	100.0%	100.0%
03-rn-flatlist-two-column-product-grid	lists	100.0%	100.0%	100.0%	100.0%
04-rn-flatlist-viewability-impression-tracking	lists	84.0%	20.0%	77.8%	100.0%
05-rn-flatlist-large-feed-window-tuning	lists	100.0%	100.0%	100.0%	100.0%
06-rn-flatlist-pull-to-refresh-feed	lists	98.0%	90.0%	100.0%	100.0%
07-rn-flatlist-to-sectionlist-conversion	lists	100.0%	100.0%	100.0%	100.0%
08-rn-flashlist-basic-list	lists	100.0%	100.0%	100.0%	100.0%
09-rn-flashlist-horizontal-rail-list-header	lists	92.5%	70.0%	100.0%	100.0%
10-rn-flashlist-horizontal-dynamic-height-header-first-item	lists	71.7%	30.0%	91.7%	100.0%
11-rn-flashlist-expandable-cards-layout-state	lists	66.7%	0.0%	0.0%	0.0%
12-rn-flashlist-chat-start-from-bottom	lists	100.0%	100.0%	100.0%	100.0%
13-rn-flashlist-masonry-photo-gallery	lists	40.0%	0.0%	0.0%	0.0%
14-rn-flatlist-to-legendlist-migration	lists	66.7%	0.0%	0.0%	0.0%
15-rn-legendlist-sticky-grouped-agenda	lists	82.0%	10.0%	50.0%	100.0%
16-rn-legendlist-section-scroll-to-location	lists	100.0%	100.0%	100.0%	100.0%
17-rn-legendlist-chat-start-from-bottom	lists	56.0%	20.0%	77.8%	100.0%
18-rn-legendlist-expandable-cards-recycling-state	lists	50.0%	0.0%	0.0%	0.0%
01-rn-nav-stack-product-details	navigation	80.0%	0.0%	0.0%	0.0%
02-rn-nav-stack-filter-default	navigation	100.0%	100.0%	100.0%	100.0%
03-rn-nav-stack-header-from-param	navigation	97.5%	90.0%	100.0%	100.0%
04-rn-nav-modal-compose-note	navigation	100.0%	100.0%	100.0%	100.0%
05-rn-nav-tabs-three-sections	navigation	100.0%	100.0%	100.0%	100.0%
06-rn-nav-home-tab-nested-details	navigation	80.0%	40.0%	97.6%	100.0%
07-rn-nav-drawer-account-help	navigation	100.0%	100.0%	100.0%	100.0%
09-rn-nav-login-replace-home	navigation	100.0%	100.0%	100.0%	100.0%
10-rn-nav-deeplink-profile-basic	navigation	100.0%	100.0%	100.0%	100.0%
12-rn-nav-unsaved-edit-confirm	navigation	92.5%	70.0%	100.0%	100.0%
13-rn-nav-focus-refresh-inbox	navigation	100.0%	100.0%	100.0%	100.0%
14-rn-nav-tabs-lazy-first-render	navigation	100.0%	100.0%	100.0%	100.0%
15-rn-nav-header-action-stateful	navigation	100.0%	100.0%	100.0%	100.0%
01-rn-stylesheet-card-variants	react native apis	100.0%	100.0%	100.0%	100.0%
02-rn-stylesheet-absolute-fill-overlay-badge	react native apis	50.0%	0.0%	0.0%	0.0%
03-rn-stylesheet-divider-hairline	react native apis	86.7%	80.0%	100.0%	100.0%
04-rn-image-network-dimensions	react native apis	100.0%	100.0%	100.0%	100.0%
05-rn-imagebackground-resize-mode	react native apis	100.0%	100.0%	100.0%	100.0%
06-rn-platform-select-small-style-delta	react native apis	100.0%	100.0%	100.0%	100.0%
07-rn-platform-os-conditional-cta	react native apis	100.0%	100.0%	100.0%	100.0%
08-rn-platform-file-extension-component-split	react native apis	100.0%	100.0%	100.0%	100.0%
09-rn-keyboardavoidingview-login-form	react native apis	100.0%	100.0%	100.0%	100.0%

Table 37: Per-eval results (score and pass@k) for gpt-5.3-codex




Figure 23: Benchmark Plot for gpt-5.3-codex.

vercel/openai/gpt-5.3-codex · Run → Category → Pass/Fail




Figure 24: Sankey per run for gpt-5.3-codex showing the distribution of evals with respect to categories and results

A.12 gpt-5.4

Model	Runs	Score	Reqs	P@1	P@5	P@10
vercel/openai/gpt-5.4	10	85.3%	2330/2730	64.4%	74.7%	75.8%

Table 38: Model summary for gpt-5.4

Category	Score	Evals	Requirements	Failed requirements
animation	71.3%	130	53	164
async state	85.6%	130	65	96
lists	86.7%	180	79	102
navigation	95.1%	130	45	23
react native apis	94.4%	90	31	15

Table 39: Results by category for gpt-5.4. Score: mean(per-eval passed/total weight) across evals and runs. Evals: count of evaluation tasks. Requirements: sum of requirement count per eval. Failed requirements: sum of failed requirement instances across runs.

Eval	Category	Score	P@1	P@5	P@10
01-rn-anim-pressable-scale-with-timing	animation	100.0%	100.0%	100.0%	100.0%
02-rn-anim-spring-toggle-card-height	animation	76.7%	40.0%	97.6%	100.0%
03-rn-anim-layout-enter-exit-list-items	animation	100.0%	100.0%	100.0%	100.0%
04-rn-anim-scroll-linked-collapsing-header	animation	100.0%	100.0%	100.0%	100.0%
05-rn-anim-pan-drag-card-with-snap-points	animation	65.0%	0.0%	0.0%	0.0%
06-rn-anim-longpress-then-pan-activation-gate	animation	85.0%	40.0%	97.6%	100.0%
07-rn-anim-pinch-pan-simultaneous-photo-canvas	animation	100.0%	100.0%	100.0%	100.0%
09-rn-anim-android-modal-gesture-root-safety	animation	50.0%	0.0%	0.0%	0.0%
10-rn-anim-worklet-thread-bridge-schedule-on-rn	animation	50.0%	0.0%	0.0%	0.0%
11-rn-rngh-exclusive-single-double-tap	animation	54.0%	0.0%	0.0%	0.0%
12-rn-rngh-scroll-pan-failure-chain	animation	80.0%	0.0%	0.0%	0.0%
13-rn-worklets-runtime-offload-and-bridge	animation	36.7%	0.0%	0.0%	0.0%
14-rn-keyboard-controller-aware-scroll-footer	animation	30.0%	0.0%	0.0%	0.0%
01-rn-rq-stable-query-key-filters	async state	44.0%	0.0%	0.0%	0.0%
02-rn-rq-dependent-query-enabled-gate	async state	58.0%	0.0%	0.0%	0.0%
03-rn-rq-mutation-invalidate-on-success	async state	60.0%	0.0%	0.0%	0.0%
04-rn-rq-optimistic-update-rollback	async state	97.1%	80.0%	100.0%	100.0%
05-rn-zustand-async-action-lifecycle	async state	100.0%	100.0%	100.0%	100.0%
06-rn-zustand-request-id-stale-response-guard	async state	92.0%	70.0%	100.0%	100.0%
07-rn-zustand-persist-async-hydration-gate	async state	82.9%	40.0%	97.6%	100.0%
09-rn-jotai-async-read-atom-suspense	async state	100.0%	100.0%	100.0%	100.0%
10-rn-jotai-loadable-async-atom-no-suspense	async state	82.5%	30.0%	91.7%	100.0%
11-rn-jotai-async-write-atom-mutation	async state	100.0%	100.0%	100.0%	100.0%
12-rn-jotai-atomwithstorage-async-rehydrate-guard	async state	96.7%	80.0%	100.0%	100.0%
13-rn-react-usedeferredvalue-search-stale-indicator	async state	100.0%	100.0%	100.0%	100.0%
14-rn-rq-react-native-appstate-focus-online-managers	async state	100.0%	100.0%	100.0%	100.0%
01-rn-flatlist-scroll-to-unread-fixed-rows	lists	100.0%	100.0%	100.0%	100.0%
02-rn-flatlist-basic-contact-list	lists	100.0%	100.0%	100.0%	100.0%
03-rn-flatlist-two-column-product-grid	lists	100.0%	100.0%	100.0%	100.0%
04-rn-flatlist-viewability-impression-tracking	lists	80.0%	0.0%	0.0%	0.0%
05-rn-flatlist-large-feed-window-tuning	lists	100.0%	100.0%	100.0%	100.0%
06-rn-flatlist-pull-to-refresh-feed	lists	98.0%	90.0%	100.0%	100.0%
07-rn-flatlist-to-sectionlist-conversion	lists	100.0%	100.0%	100.0%	100.0%
08-rn-flashlist-basic-list	lists	100.0%	100.0%	100.0%	100.0%
09-rn-flashlist-horizontal-rail-list-header	lists	100.0%	100.0%	100.0%	100.0%
10-rn-flashlist-horizontal-dynamic-height-header-first-item	lists	90.0%	40.0%	97.6%	100.0%
11-rn-flashlist-expandable-cards-layout-state	lists	66.7%	0.0%	0.0%	0.0%
12-rn-flashlist-chat-start-from-bottom	lists	100.0%	100.0%	100.0%	100.0%
13-rn-flashlist-masonry-photo-gallery	lists	52.0%	20.0%	77.8%	100.0%
14-rn-flatlist-to-legendlist-migration	lists	66.7%	0.0%	0.0%	0.0%
15-rn-legendlist-sticky-grouped-agenda	lists	82.0%	30.0%	91.7%	100.0%
16-rn-legendlist-section-scroll-to-location	lists	100.0%	100.0%	100.0%	100.0%
17-rn-legendlist-chat-start-from-bottom	lists	76.0%	20.0%	77.8%	100.0%
18-rn-legendlist-expandable-cards-recycling-state	lists	50.0%	0.0%	0.0%	0.0%
01-rn-nav-stack-product-details	navigation	94.0%	70.0%	100.0%	100.0%
02-rn-nav-stack-filter-default	navigation	100.0%	100.0%	100.0%	100.0%
03-rn-nav-stack-header-from-param	navigation	92.5%	70.0%	100.0%	100.0%
04-rn-nav-modal-compose-note	navigation	100.0%	100.0%	100.0%	100.0%
05-rn-nav-tabs-three-sections	navigation	100.0%	100.0%	100.0%	100.0%
06-rn-nav-home-tab-nested-details	navigation	80.0%	40.0%	97.6%	100.0%
07-rn-nav-drawer-account-help	navigation	100.0%	100.0%	100.0%	100.0%
09-rn-nav-login-replace-home	navigation	100.0%	100.0%	100.0%	100.0%
10-rn-nav-deeplink-profile-basic	navigation	100.0%	100.0%	100.0%	100.0%
12-rn-nav-unsaved-edit-confirm	navigation	75.0%	0.0%	0.0%	0.0%
13-rn-nav-focus-refresh-inbox	navigation	100.0%	100.0%	100.0%	100.0%
14-rn-nav-tabs-lazy-first-render	navigation	100.0%	100.0%	100.0%	100.0%
15-rn-nav-header-action-stateful	navigation	95.0%	90.0%	100.0%	100.0%
01-rn-stylesheet-card-variants	react native apis	100.0%	100.0%	100.0%	100.0%
02-rn-stylesheet-absolute-fill-overlay-badge	react native apis	50.0%	0.0%	0.0%	0.0%
03-rn-stylesheet-divider-hairline	react native apis	100.0%	100.0%	100.0%	100.0%
04-rn-image-network-dimensions	react native apis	100.0%	100.0%	100.0%	100.0%
05-rn-imagebackground-resize-mode	react native apis	100.0%	100.0%	100.0%	100.0%
06-rn-platform-select-small-style-delta	react native apis	100.0%	100.0%	100.0%	100.0%
07-rn-platform-os-conditional-cta	react native apis	100.0%	100.0%	100.0%	100.0%
08-rn-platform-file-extension-component-split	react native apis	100.0%	100.0%	100.0%	100.0%
09-rn-keyboardavoidingview-login-form	react native apis	100.0%	100.0%	100.0%	100.0%

Table 40: Per-eval results (score and pass@k) for gpt-5.4




Figure 25: Benchmark Plot for gpt-5.4.

vercel/openai/gpt-5.4 · Run → Category → Pass/Fail




Figure 26: Sankey per run for gpt-5.4 showing the distribution of evals with respect to categories and results

A.13 gpt-oss-120b

Model	Runs	Score	Reqs	P@1	P@5	P@10
cloudflare-ai-gateway/workers-ai/@cf/openai/gpt-os	10	71.6%	1653/2308	46.0%	54.3%	50.4%

Table 41: Model summary for gpt-oss-120b

Category	Score	Evals	Requirements	Failed requirements
animation	52.6%	110	53	239
async state	67.1%	129	65	204
lists	73.0%	122	79	144
navigation	88.5%	109	45	51
react native apis	93.7%	89	31	17

Table 42: Results by category for gpt-oss-120b. Score: mean(per-eval passed/total weight) across evals and runs. Evals: count of evaluation tasks. Requirements: sum of requirement count per eval. Failed requirements: sum of failed requirement instances across runs.

Eval	Category	Score	P@1	P@5	P@10
01-rn-anim-pressable-scale-with-timing	animation	100.0%	100.0%	100.0%	100.0%
02-rn-anim-spring-toggle-card-height	animation	60.0%	0.0%	0.0%	0.0%
03-rn-anim-layout-enter-exit-list-items	animation	100.0%	100.0%	100.0%	100.0%
04-rn-anim-scroll-linked-collapsing-header	animation	100.0%	100.0%	100.0%	100.0%
05-rn-anim-pan-drag-card-with-snap-points	animation	25.0%	0.0%	0.0%	0.0%
06-rn-anim-longpress-then-pan-activation-gate	animation	45.0%	0.0%	0.0%	0.0%
07-rn-anim-pinch-pan-simultaneous-photo-canvas	animation	96.0%	80.0%	100.0%	80.0%
09-rn-anim-android-modal-gesture-root-safety	animation	16.7%	0.0%	0.0%	0.0%
10-rn-anim-worklet-thread-bridge-schedule-on-rn	animation	47.5%	0.0%	0.0%	0.0%
11-rn-rngh-exclusive-single-double-tap	animation	40.0%	0.0%	0.0%	0.0%
12-rn-rngh-scroll-pan-failure-chain	animation	12.0%	0.0%	0.0%	0.0%
13-rn-worklets-runtime-offload-and-bridge	animation	0.0%	0.0%	0.0%	0.0%
14-rn-keyboard-controller-aware-scroll-footer	animation	3.3%	0.0%	0.0%	0.0%
01-rn-rq-stable-query-key-filters	async state	20.0%	0.0%	0.0%	0.0%
02-rn-rq-dependent-query-enabled-gate	async state	40.0%	0.0%	0.0%	0.0%
03-rn-rq-mutation-invalidate-on-success	async state	36.0%	0.0%	0.0%	0.0%
04-rn-rq-optimistic-update-rollback	async state	84.3%	40.0%	97.6%	100.0%
05-rn-zustand-async-action-lifecycle	async state	100.0%	100.0%	100.0%	100.0%
06-rn-zustand-request-id-stale-response-guard	async state	88.0%	40.0%	97.6%	100.0%
07-rn-zustand-persist-async-hydration-gate	async state	54.3%	0.0%	0.0%	0.0%
09-rn-jotai-async-read-atom-suspense	async state	100.0%	100.0%	100.0%	100.0%
10-rn-jotai-loadable-async-atom-no-suspense	async state	27.5%	0.0%	0.0%	0.0%
11-rn-jotai-async-write-atom-mutation	async state	100.0%	100.0%	100.0%	100.0%
12-rn-jotai-atomwithstorage-async-rehydrate-guard	async state	96.7%	80.0%	100.0%	100.0%
13-rn-react-usedeferredvalue-search-stale-indicator	async state	94.4%	88.9%	100.0%	88.9%
14-rn-rq-react-native-appstate-focus-online-managers	async state	33.3%	0.0%	0.0%	0.0%
01-rn-flatlist-scroll-to-unread-fixed-rows	lists	74.3%	14.3%	71.4%	14.3%
02-rn-flatlist-basic-contact-list	lists	100.0%	100.0%	100.0%	100.0%
03-rn-flatlist-two-column-product-grid	lists	100.0%	100.0%	100.0%	100.0%
04-rn-flatlist-viewability-impression-tracking	lists	60.0%	0.0%	0.0%	0.0%
05-rn-flatlist-large-feed-window-tuning	lists	76.0%	60.0%	100.0%	60.0%
06-rn-flatlist-pull-to-refresh-feed	lists	80.0%	0.0%	0.0%	0.0%
07-rn-flatlist-to-sectionlist-conversion	lists	100.0%	100.0%	100.0%	100.0%
08-rn-flashlist-basic-list	lists	100.0%	100.0%	100.0%	100.0%
09-rn-flashlist-horizontal-rail-list-header	lists	75.0%	10.0%	50.0%	100.0%
10-rn-flashlist-horizontal-dynamic-height-header-first-item	lists	55.6%	0.0%	0.0%	0.0%
11-rn-flashlist-expandable-cards-layout-state	lists	66.7%	0.0%	0.0%	0.0%
12-rn-flashlist-chat-start-from-bottom	lists	50.0%	0.0%	0.0%	0.0%
13-rn-flashlist-masonry-photo-gallery	lists	76.0%	0.0%	0.0%	0.0%
14-rn-flatlist-to-legendlist-migration	lists	0.0%	0.0%	0.0%	0.0%
15-rn-legendlist-sticky-grouped-agenda	lists	80.0%	0.0%	0.0%	0.0%
16-rn-legendlist-section-scroll-to-location	lists	100.0%	100.0%	100.0%	100.0%
17-rn-legendlist-chat-start-from-bottom	lists	40.0%	0.0%	0.0%	0.0%
18-rn-legendlist-expandable-cards-recycling-state	lists	50.0%	0.0%	0.0%	0.0%
01-rn-nav-stack-product-details	navigation	73.3%	0.0%	0.0%	0.0%
02-rn-nav-stack-filter-default	navigation	100.0%	100.0%	100.0%	100.0%
03-rn-nav-stack-header-from-param	navigation	80.6%	22.2%	83.3%	22.2%
04-rn-nav-modal-compose-note	navigation	95.0%	80.0%	100.0%	100.0%
05-rn-nav-tabs-three-sections	navigation	72.2%	16.7%	83.3%	16.7%
06-rn-nav-home-tab-nested-details	navigation	96.3%	88.9%	100.0%	88.9%
07-rn-nav-drawer-account-help	navigation	87.5%	66.7%	100.0%	66.7%
09-rn-nav-login-replace-home	navigation	76.0%	0.0%	0.0%	0.0%
10-rn-nav-deeplink-profile-basic	navigation	100.0%	100.0%	100.0%	100.0%
12-rn-nav-unsaved-edit-confirm	navigation	62.5%	0.0%	0.0%	0.0%
13-rn-nav-focus-refresh-inbox	navigation	100.0%	100.0%	100.0%	100.0%
14-rn-nav-tabs-lazy-first-render	navigation	100.0%	100.0%	100.0%	100.0%
15-rn-nav-header-action-stateful	navigation	95.0%	90.0%	100.0%	100.0%
01-rn-stylesheet-card-variants	react native apis	100.0%	100.0%	100.0%	100.0%
02-rn-stylesheet-absolute-fill-overlay-badge	react native apis	60.0%	0.0%	0.0%	0.0%
03-rn-stylesheet-divider-hairline	react native apis	90.0%	80.0%	100.0%	100.0%
04-rn-image-network-dimensions	react native apis	100.0%	100.0%	100.0%	100.0%
05-rn-imagebackground-resize-mode	react native apis	97.5%	90.0%	100.0%	100.0%
06-rn-platform-select-small-style-delta	react native apis	100.0%	100.0%	100.0%	100.0%
07-rn-platform-os-conditional-cta	react native apis	100.0%	100.0%	100.0%	100.0%
08-rn-platform-file-extension-component-split	react native apis	100.0%	100.0%	100.0%	100.0%
09-rn-keyboardavoidingview-login-form	react native apis	96.3%	88.9%	100.0%	88.9%

Table 43: Per-eval results (score and pass@k) for gpt-oss-120b

Benchmark (10 runs) · cloudflare-ai-gateway/workers-ai/@cf/openai/gpt-oss-120b · score 71.6% (1653/2308 reqs) · pass@1/5/10: 46.0% / 54.3% / 50.4% · cost \$2.56



Figure 27: Benchmark Plot for gpt-oss-120b.

cloudflare-ai-gateway/workers-ai/@cf/openai/gpt-oss-120b · Run → Category → Pass/Fail




Figure 28: Sankey per run for gpt-oss-120b showing the distribution of evals with respect to categories and results

A.14 gpt-oss-20b

Model	Runs	Score	Reqs	P@1	P@5	P@10
cloudflare-ai-gateway/workers-ai/@cf/openai/gpt-os	10	71.0%	1761/2480	43.4%	56.2%	55.4%

Table 44: Model summary for gpt-oss-20b

Category	Score	Evals	Requirements	Failed requirements
animation	51.7%	107	53	232
async state	65.8%	128	65	212
lists	75.0%	175	79	193
navigation	84.2%	94	45	61
react native apis	92.9%	88	31	21

Table 45: Results by category for gpt-oss-20b. Score: mean(per-eval passed/total weight) across evals and runs. Evals: count of evaluation tasks. Requirements: sum of requirement count per eval. Failed requirements: sum of failed requirement instances across runs.

Eval	Category	Score	P@1	P@5	P@10
01-rn-anim-pressable-scale-with-timing	animation	100.0%	100.0%	100.0%	100.0%
02-rn-anim-spring-toggle-card-height	animation	53.3%	0.0%	0.0%	0.0%
03-rn-anim-layout-enter-exit-list-items	animation	60.0%	40.0%	97.6%	100.0%
04-rn-anim-scroll-linked-collapsing-header	animation	100.0%	100.0%	100.0%	100.0%
05-rn-anim-pan-drag-card-with-snap-points	animation	22.5%	0.0%	0.0%	0.0%
06-rn-anim-longpress-then-pan-activation-gate	animation	47.5%	0.0%	0.0%	0.0%
07-rn-anim-pinch-pan-simultaneous-photo-canvas	animation	88.0%	40.0%	97.6%	100.0%
09-rn-anim-android-modal-gesture-root-safety	animation	21.7%	0.0%	0.0%	0.0%
10-rn-anim-worklet-thread-bridge-schedule-on-rn	animation	50.0%	0.0%	0.0%	0.0%
11-rn-rngh-exclusive-single-double-tap	animation	10.0%	0.0%	0.0%	0.0%
12-rn-rngh-scroll-pan-failure-chain	animation	22.0%	0.0%	0.0%	0.0%
13-rn-worklets-runtime-offload-and-bridge	animation	33.3%	0.0%	0.0%	0.0%
14-rn-keyboard-controller-aware-scroll-footer	animation	0.0%	0.0%	0.0%	0.0%
01-rn-rq-stable-query-key-filters	async state	18.0%	0.0%	0.0%	0.0%
02-rn-rq-dependent-query-enabled-gate	async state	48.0%	0.0%	0.0%	0.0%
03-rn-rq-mutation-invalidate-on-success	async state	38.0%	0.0%	0.0%	0.0%
04-rn-rq-optimistic-update-rollback	async state	81.4%	10.0%	50.0%	100.0%
05-rn-zustand-async-action-lifecycle	async state	100.0%	100.0%	100.0%	100.0%
06-rn-zustand-request-id-stale-response-guard	async state	82.0%	40.0%	97.6%	100.0%
07-rn-zustand-persist-async-hydration-gate	async state	42.9%	0.0%	0.0%	0.0%
09-rn-jotai-async-read-atom-suspense	async state	97.5%	90.0%	100.0%	100.0%
10-rn-jotai-loadable-async-atom-no-suspense	async state	41.7%	11.1%	55.6%	11.1%
11-rn-jotai-async-write-atom-mutation	async state	94.0%	90.0%	100.0%	100.0%
12-rn-jotai-atomwithstorage-async-rehydrate-guard	async state	98.3%	90.0%	100.0%	100.0%
13-rn-react-usedeferredvalue-search-stale-indicator	async state	75.0%	40.0%	97.6%	100.0%
14-rn-rq-react-native-appstate-focus-online-managers	async state	33.3%	0.0%	0.0%	0.0%
01-rn-flatlist-scroll-to-unread-fixed-rows	lists	60.0%	30.0%	91.7%	100.0%
02-rn-flatlist-basic-contact-list	lists	100.0%	100.0%	100.0%	100.0%
03-rn-flatlist-two-column-product-grid	lists	100.0%	100.0%	100.0%	100.0%
04-rn-flatlist-viewability-impression-tracking	lists	80.0%	0.0%	0.0%	0.0%
05-rn-flatlist-large-feed-window-tuning	lists	92.0%	60.0%	100.0%	100.0%
06-rn-flatlist-pull-to-refresh-feed	lists	92.0%	70.0%	100.0%	100.0%
07-rn-flatlist-to-sectionlist-conversion	lists	100.0%	100.0%	100.0%	100.0%
08-rn-flashlist-basic-list	lists	100.0%	100.0%	100.0%	100.0%
09-rn-flashlist-horizontal-rail-list-header	lists	92.5%	70.0%	100.0%	100.0%
10-rn-flashlist-horizontal-dynamic-height-header-first-item	lists	66.7%	0.0%	0.0%	0.0%
11-rn-flashlist-expandable-cards-layout-state	lists	66.7%	0.0%	0.0%	0.0%
12-rn-flashlist-chat-start-from-bottom	lists	46.0%	0.0%	0.0%	0.0%
13-rn-flashlist-masonry-photo-gallery	lists	76.0%	10.0%	50.0%	100.0%
14-rn-flatlist-to-legendlist-migration	lists	0.0%	0.0%	0.0%	0.0%
15-rn-legendlist-sticky-grouped-agenda	lists	82.0%	10.0%	50.0%	100.0%
16-rn-legendlist-section-scroll-to-location	lists	100.0%	100.0%	100.0%	100.0%
17-rn-legendlist-chat-start-from-bottom	lists	42.2%	0.0%	0.0%	0.0%
18-rn-legendlist-expandable-cards-recycling-state	lists	25.0%	0.0%	0.0%	0.0%
01-rn-nav-stack-product-details	navigation	80.0%	11.1%	55.6%	11.1%
02-rn-nav-stack-filter-default	navigation	100.0%	100.0%	100.0%	100.0%
03-rn-nav-stack-header-from-param	navigation	80.0%	20.0%	100.0%	20.0%
04-rn-nav-modal-compose-note	navigation	97.5%	90.0%	100.0%	100.0%
05-rn-nav-tabs-three-sections	navigation	66.7%	0.0%	0.0%	0.0%
06-rn-nav-home-tab-nested-details	navigation	66.7%	14.3%	71.4%	14.3%
07-rn-nav-drawer-account-help	navigation	50.0%	0.0%	0.0%	0.0%
09-rn-nav-login-replace-home	navigation	64.4%	0.0%	0.0%	0.0%
10-rn-nav-deeplink-profile-basic	navigation	100.0%	100.0%	100.0%	100.0%
12-rn-nav-unsaved-edit-confirm	navigation	52.5%	0.0%	0.0%	0.0%
13-rn-nav-focus-refresh-inbox	navigation	100.0%	100.0%	100.0%	100.0%
14-rn-nav-tabs-lazy-first-render	navigation	100.0%	100.0%	100.0%	100.0%
15-rn-nav-header-action-stateful	navigation	100.0%	100.0%	100.0%	100.0%
01-rn-stylesheet-card-variants	react native apis	93.3%	90.0%	100.0%	100.0%
02-rn-stylesheet-absolute-fill-overlay-badge	react native apis	66.7%	0.0%	0.0%	0.0%
03-rn-stylesheet-divider-hairline	react native apis	100.0%	100.0%	100.0%	100.0%
04-rn-image-network-dimensions	react native apis	100.0%	100.0%	100.0%	100.0%
05-rn-imagebackground-resize-mode	react native apis	77.5%	40.0%	97.6%	100.0%
06-rn-platform-select-small-style-delta	react native apis	100.0%	100.0%	100.0%	100.0%
07-rn-platform-os-conditional-cta	react native apis	100.0%	100.0%	100.0%	100.0%
08-rn-platform-file-extension-component-split	react native apis	100.0%	100.0%	100.0%	100.0%
09-rn-keyboardavoidingview-login-form	react native apis	100.0%	100.0%	100.0%	100.0%

Table 46: Per-eval results (score and pass@k) for gpt-oss-20b




Figure 29: Benchmark Plot for gpt-oss-20b.

cloudflare-ai-gateway/workers-ai/@cf/openai/gpt-oss-20b · Run → Category → Pass/Fail




Figure 30: Sankey per run for gpt-oss-20b showing the distribution of evals with respect to categories and results

A.15 grok-4

Model	Runs	Score	Reqs	P@1	P@5	P@10
vercel/xai/grok-4	10	72.6%	1980/2726	44.7%	55.1%	56.1%

Table 47: Model summary for grok-4

Category	Score	Evals	Requirements	Failed requirements
animation	60.7%	130	53	226
async state	70.1%	130	65	187
lists	69.6%	180	79	237
navigation	84.9%	130	45	75
react native apis	92.1%	89	31	21

Table 48: Results by category for grok-4. Score: mean(per-eval passed/total weight) across evals and runs. Evals: count of evaluation tasks. Requirements: sum of requirement count per eval. Failed requirements: sum of failed requirement instances across runs.

Eval	Category	Score	P@1	P@5	P@10
01-rn-anim-pressable-scale-with-timing	animation	100.0%	100.0%	100.0%	100.0%
02-rn-anim-spring-toggle-card-height	animation	63.3%	0.0%	0.0%	0.0%
03-rn-anim-layout-enter-exit-list-items	animation	93.3%	90.0%	100.0%	100.0%
04-rn-anim-scroll-linked-collapsing-header	animation	96.7%	90.0%	100.0%	100.0%
05-rn-anim-pan-drag-card-with-snap-points	animation	45.0%	0.0%	0.0%	0.0%
06-rn-anim-longpress-then-pan-activation-gate	animation	75.0%	50.0%	99.6%	100.0%
07-rn-anim-pinch-pan-simultaneous-photo-canvas	animation	100.0%	100.0%	100.0%	100.0%
09-rn-anim-android-modal-gesture-root-safety	animation	26.7%	0.0%	0.0%	0.0%
10-rn-anim-worklet-thread-bridge-schedule-on-rn	animation	45.0%	0.0%	0.0%	0.0%
11-rn-rngh-exclusive-single-double-tap	animation	42.0%	0.0%	0.0%	0.0%
12-rn-rngh-scroll-pan-failure-chain	animation	54.0%	0.0%	0.0%	0.0%
13-rn-worklets-runtime-offload-and-bridge	animation	30.0%	0.0%	0.0%	0.0%
14-rn-keyboard-controller-aware-scroll-footer	animation	18.0%	0.0%	0.0%	0.0%
01-rn-rq-stable-query-key-filters	async state	40.0%	0.0%	0.0%	0.0%
02-rn-rq-dependent-query-enabled-gate	async state	40.0%	0.0%	0.0%	0.0%
03-rn-rq-mutation-invalidate-on-success	async state	50.0%	0.0%	0.0%	0.0%
04-rn-rq-optimistic-update-rollback	async state	82.9%	20.0%	77.8%	100.0%
05-rn-zustand-async-action-lifecycle	async state	100.0%	100.0%	100.0%	100.0%
06-rn-zustand-request-id-stale-response-guard	async state	68.0%	0.0%	0.0%	0.0%
07-rn-zustand-persist-async-hydration-gate	async state	62.9%	0.0%	0.0%	0.0%
09-rn-jotai-async-read-atom-suspense	async state	97.5%	90.0%	100.0%	100.0%
10-rn-jotai-loadable-async-atom-no-suspense	async state	50.0%	0.0%	0.0%	0.0%
11-rn-jotai-async-write-atom-mutation	async state	100.0%	100.0%	100.0%	100.0%
12-rn-jotai-atomwithstorage-async-rehydrate-guard	async state	95.0%	70.0%	100.0%	100.0%
13-rn-react-usedeferredvalue-search-stale-indicator	async state	85.0%	80.0%	100.0%	100.0%
14-rn-rq-react-native-appstate-focus-online-managers	async state	40.0%	0.0%	0.0%	0.0%
01-rn-flatlist-scroll-to-unread-fixed-rows	lists	72.0%	30.0%	91.7%	100.0%
02-rn-flatlist-basic-contact-list	lists	100.0%	100.0%	100.0%	100.0%
03-rn-flatlist-two-column-product-grid	lists	100.0%	100.0%	100.0%	100.0%
04-rn-flatlist-viewability-impression-tracking	lists	80.0%	0.0%	0.0%	0.0%
05-rn-flatlist-large-feed-window-tuning	lists	86.0%	30.0%	91.7%	100.0%
06-rn-flatlist-pull-to-refresh-feed	lists	86.0%	30.0%	91.7%	100.0%
07-rn-flatlist-to-sectionlist-conversion	lists	100.0%	100.0%	100.0%	100.0%
08-rn-flashlist-basic-list	lists	86.7%	60.0%	100.0%	100.0%
09-rn-flashlist-horizontal-rail-list-header	lists	87.5%	50.0%	99.6%	100.0%
10-rn-flashlist-horizontal-dynamic-height-header-first-item	lists	60.0%	0.0%	0.0%	0.0%
11-rn-flashlist-expandable-cards-layout-state	lists	33.3%	0.0%	0.0%	0.0%
12-rn-flashlist-chat-start-from-bottom	lists	40.0%	0.0%	0.0%	0.0%
13-rn-flashlist-masonry-photo-gallery	lists	58.0%	0.0%	0.0%	0.0%
14-rn-flatlist-to-legendlist-migration	lists	20.0%	0.0%	0.0%	0.0%
15-rn-legendlist-sticky-grouped-agenda	lists	80.0%	0.0%	0.0%	0.0%
16-rn-legendlist-section-scroll-to-location	lists	100.0%	100.0%	100.0%	100.0%
17-rn-legendlist-chat-start-from-bottom	lists	38.0%	0.0%	0.0%	0.0%
18-rn-legendlist-expandable-cards-recycling-state	lists	25.0%	0.0%	0.0%	0.0%
01-rn-nav-stack-product-details	navigation	78.0%	0.0%	0.0%	0.0%
02-rn-nav-stack-filter-default	navigation	96.0%	80.0%	100.0%	100.0%
03-rn-nav-stack-header-from-param	navigation	90.0%	60.0%	100.0%	100.0%
04-rn-nav-modal-compose-note	navigation	100.0%	100.0%	100.0%	100.0%
05-rn-nav-tabs-three-sections	navigation	76.7%	30.0%	91.7%	100.0%
06-rn-nav-home-tab-nested-details	navigation	56.7%	0.0%	0.0%	0.0%
07-rn-nav-drawer-account-help	navigation	67.5%	30.0%	91.7%	100.0%
09-rn-nav-login-replace-home	navigation	74.0%	0.0%	0.0%	0.0%
10-rn-nav-deeplink-profile-basic	navigation	100.0%	100.0%	100.0%	100.0%
12-rn-nav-unsaved-edit-confirm	navigation	75.0%	0.0%	0.0%	0.0%
13-rn-nav-focus-refresh-inbox	navigation	100.0%	100.0%	100.0%	100.0%
14-rn-nav-tabs-lazy-first-render	navigation	100.0%	100.0%	100.0%	100.0%
15-rn-nav-header-action-stateful	navigation	90.0%	80.0%	100.0%	100.0%
01-rn-stylesheet-card-variants	react native apis	100.0%	100.0%	100.0%	100.0%
02-rn-stylesheet-absolute-fill-overlay-badge	react native apis	36.7%	0.0%	0.0%	0.0%
03-rn-stylesheet-divider-hairline	react native apis	93.3%	80.0%	100.0%	100.0%
04-rn-image-network-dimensions	react native apis	100.0%	100.0%	100.0%	100.0%
05-rn-imagebackground-resize-mode	react native apis	100.0%	100.0%	100.0%	100.0%
06-rn-platform-select-small-style-delta	react native apis	100.0%	100.0%	100.0%	100.0%
07-rn-platform-os-conditional-cta	react native apis	100.0%	100.0%	100.0%	100.0%
08-rn-platform-file-extension-component-split	react native apis	100.0%	100.0%	100.0%	100.0%
09-rn-keyboardavoidingview-login-form	react native apis	100.0%	100.0%	100.0%	100.0%

Table 49: Per-eval results (score and pass@k) for grok-4




Figure 31: Benchmark Plot for grok-4.

vercel/xai/grok-4 · Run → Category → Pass/Fail




Figure 32: Sankey per run for grok-4 showing the distribution of evals with respect to categories and results

A.16 kimi-k2.5

Model	Runs	Score	Reqs	P@1	P@5	P@10
vercel/moonshotai/kimi-k2.5	10	77.2%	2107/2730	49.4%	65.6%	69.7%

Table 50: Model summary for kimi-k2.5

Category	Score	Evals	Requirements	Failed requirements
animation	61.0%	130	53	220
async state	76.0%	130	65	148
lists	77.7%	180	79	181
navigation	91.9%	130	45	41
react native apis	88.3%	90	31	33

Table 51: Results by category for kimi-k2.5. Score: mean(per-eval passed/total weight) across evals and runs. Evals: count of evaluation tasks. Requirements: sum of requirement count per eval. Failed requirements: sum of failed requirement instances across runs.

Eval	Category	Score	P@1	P@5	P@10
01-rn-anim-pressable-scale-with-timing	animation	100.0%	100.0%	100.0%	100.0%
02-rn-anim-spring-toggle-card-height	animation	66.7%	30.0%	91.7%	100.0%
03-rn-anim-layout-enter-exit-list-items	animation	100.0%	100.0%	100.0%	100.0%
04-rn-anim-scroll-linked-collapsing-header	animation	73.3%	40.0%	97.6%	100.0%
05-rn-anim-pan-drag-card-with-snap-points	animation	55.0%	0.0%	0.0%	0.0%
06-rn-anim-longpress-then-pan-activation-gate	animation	72.5%	40.0%	97.6%	100.0%
07-rn-anim-pinch-pan-simultaneous-photo-canvas	animation	100.0%	100.0%	100.0%	100.0%
09-rn-anim-android-modal-gesture-root-safety	animation	30.0%	0.0%	0.0%	0.0%
10-rn-anim-worklet-thread-bridge-schedule-on-rn	animation	20.0%	0.0%	0.0%	0.0%
11-rn-rngh-exclusive-single-double-tap	animation	38.0%	0.0%	0.0%	0.0%
12-rn-rngh-scroll-pan-failure-chain	animation	70.0%	0.0%	0.0%	0.0%
13-rn-worklets-runtime-offload-and-bridge	animation	33.3%	0.0%	0.0%	0.0%
14-rn-keyboard-controller-aware-scroll-footer	animation	34.0%	0.0%	0.0%	0.0%
01-rn-rq-stable-query-key-filters	async state	40.0%	0.0%	0.0%	0.0%
02-rn-rq-dependent-query-enabled-gate	async state	46.0%	0.0%	0.0%	0.0%
03-rn-rq-mutation-invalidate-on-success	async state	68.0%	0.0%	0.0%	0.0%
04-rn-rq-optimistic-update-rollback	async state	98.6%	90.0%	100.0%	100.0%
05-rn-zustand-async-action-lifecycle	async state	100.0%	100.0%	100.0%	100.0%
06-rn-zustand-request-id-stale-response-guard	async state	68.0%	20.0%	77.8%	100.0%
07-rn-zustand-persist-async-hydration-gate	async state	72.9%	30.0%	91.7%	100.0%
09-rn-jotai-async-read-atom-suspense	async state	97.5%	90.0%	100.0%	100.0%
10-rn-jotai-loadable-async-atom-no-suspense	async state	50.0%	0.0%	0.0%	0.0%
11-rn-jotai-async-write-atom-mutation	async state	100.0%	100.0%	100.0%	100.0%
12-rn-jotai-atomwithstorage-async-rehydrate-guard	async state	93.3%	60.0%	100.0%	100.0%
13-rn-react-usedeferredvalue-search-stale-indicator	async state	100.0%	100.0%	100.0%	100.0%
14-rn-rq-react-native-appstate-focus-online-managers	async state	53.3%	0.0%	0.0%	0.0%
01-rn-flatlist-scroll-to-unread-fixed-rows	lists	90.0%	70.0%	100.0%	100.0%
02-rn-flatlist-basic-contact-list	lists	100.0%	100.0%	100.0%	100.0%
03-rn-flatlist-two-column-product-grid	lists	100.0%	100.0%	100.0%	100.0%
04-rn-flatlist-viewability-impression-tracking	lists	80.0%	0.0%	0.0%	0.0%
05-rn-flatlist-large-feed-window-tuning	lists	86.0%	30.0%	91.7%	100.0%
06-rn-flatlist-pull-to-refresh-feed	lists	84.0%	20.0%	77.8%	100.0%
07-rn-flatlist-to-sectionlist-conversion	lists	86.0%	30.0%	91.7%	100.0%
08-rn-flashlist-basic-list	lists	100.0%	100.0%	100.0%	100.0%
09-rn-flashlist-horizontal-rail-list-header	lists	80.0%	20.0%	77.8%	100.0%
10-rn-flashlist-horizontal-dynamic-height-header-first-item	lists	83.3%	0.0%	0.0%	0.0%
11-rn-flashlist-expandable-cards-layout-state	lists	60.0%	0.0%	0.0%	0.0%
12-rn-flashlist-chat-start-from-bottom	lists	54.0%	0.0%	0.0%	0.0%
13-rn-flashlist-masonry-photo-gallery	lists	60.0%	0.0%	0.0%	0.0%
14-rn-flatlist-to-legendlist-migration	lists	70.0%	10.0%	50.0%	100.0%
15-rn-legendlist-sticky-grouped-agenda	lists	82.0%	20.0%	77.8%	100.0%
16-rn-legendlist-section-scroll-to-location	lists	100.0%	100.0%	100.0%	100.0%
17-rn-legendlist-chat-start-from-bottom	lists	40.0%	0.0%	0.0%	0.0%
18-rn-legendlist-expandable-cards-recycling-state	lists	42.5%	0.0%	0.0%	0.0%
01-rn-nav-stack-product-details	navigation	78.0%	20.0%	77.8%	100.0%
02-rn-nav-stack-filter-default	navigation	100.0%	100.0%	100.0%	100.0%
03-rn-nav-stack-header-from-param	navigation	87.5%	50.0%	99.6%	100.0%
04-rn-nav-modal-compose-note	navigation	100.0%	100.0%	100.0%	100.0%
05-rn-nav-tabs-three-sections	navigation	100.0%	100.0%	100.0%	100.0%
06-rn-nav-home-tab-nested-details	navigation	86.7%	60.0%	100.0%	100.0%
07-rn-nav-drawer-account-help	navigation	100.0%	100.0%	100.0%	100.0%
09-rn-nav-login-replace-home	navigation	72.0%	10.0%	50.0%	100.0%
10-rn-nav-deeplink-profile-basic	navigation	95.0%	90.0%	100.0%	100.0%
12-rn-nav-unsaved-edit-confirm	navigation	95.0%	80.0%	100.0%	100.0%
13-rn-nav-focus-refresh-inbox	navigation	100.0%	100.0%	100.0%	100.0%
14-rn-nav-tabs-lazy-first-render	navigation	100.0%	100.0%	100.0%	100.0%
15-rn-nav-header-action-stateful	navigation	80.0%	60.0%	100.0%	100.0%
01-rn-stylesheet-card-variants	react native apis	100.0%	100.0%	100.0%	100.0%
02-rn-stylesheet-absolute-fill-overlay-badge	react native apis	40.0%	0.0%	0.0%	0.0%
03-rn-stylesheet-divider-hairline	react native apis	70.0%	20.0%	77.8%	100.0%
04-rn-image-network-dimensions	react native apis	100.0%	100.0%	100.0%	100.0%
05-rn-imagebackground-resize-mode	react native apis	85.0%	70.0%	100.0%	100.0%
06-rn-platform-select-small-style-delta	react native apis	100.0%	100.0%	100.0%	100.0%
07-rn-platform-os-conditional-cta	react native apis	100.0%	100.0%	100.0%	100.0%
08-rn-platform-file-extension-component-split	react native apis	100.0%	100.0%	100.0%	100.0%
09-rn-keyboardavoidingview-login-form	react native apis	100.0%	100.0%	100.0%	100.0%

Table 52: Per-eval results (score and pass@k) for kimi-k2.5




Figure 33: Benchmark Plot for kimi-k2.5.

vercel/moonshotai/kimi-k2.5 · Run → Category → Pass/Fail




Figure 34: Sankey per run for kimi-k2.5 showing the distribution of evals with respect to categories and results

A.17 minimax-m2.7

Model	Runs	Score	Reqs	P@1	P@5	P@10
vercel/minimax/minimax-m2.7	10	71.4%	1948/2730	43.3%	56.1%	59.1%

Table 53: Model summary for minimax-m2.7

Category	Score	Evals	Requirements	Failed requirements
animation	58.9%	130	53	234
async state	66.2%	130	65	211
lists	68.5%	180	79	244
navigation	88.5%	130	45	61
react native apis	88.2%	90	31	32

Table 54: Results by category for minimax-m2.7. Score: mean(per-eval passed/total weight) across evals and runs. Evals: count of evaluation tasks. Requirements: sum of requirement count per eval. Failed requirements: sum of failed requirement instances across runs.

Eval	Category	Score	P@1	P@5	P@10
01-rn-anim-pressable-scale-with-timing	animation	100.0%	100.0%	100.0%	100.0%
02-rn-anim-spring-toggle-card-height	animation	73.3%	20.0%	77.8%	100.0%
03-rn-anim-layout-enter-exit-list-items	animation	93.3%	90.0%	100.0%	100.0%
04-rn-anim-scroll-linked-collapsing-header	animation	80.0%	70.0%	100.0%	100.0%
05-rn-anim-pan-drag-card-with-snap-points	animation	47.5%	0.0%	0.0%	0.0%
06-rn-anim-longpress-then-pan-activation-gate	animation	62.5%	30.0%	91.7%	100.0%
07-rn-anim-pinch-pan-simultaneous-photo-canvas	animation	90.0%	80.0%	100.0%	100.0%
09-rn-anim-android-modal-gesture-root-safety	animation	48.3%	0.0%	0.0%	0.0%
10-rn-anim-worklet-thread-bridge-schedule-on-rn	animation	45.0%	0.0%	0.0%	0.0%
11-rn-rngh-exclusive-single-double-tap	animation	36.0%	0.0%	0.0%	0.0%
12-rn-rngh-scroll-pan-failure-chain	animation	40.0%	0.0%	0.0%	0.0%
13-rn-worklets-runtime-offload-and-bridge	animation	33.3%	0.0%	0.0%	0.0%
14-rn-keyboard-controller-aware-scroll-footer	animation	16.0%	0.0%	0.0%	0.0%
01-rn-rq-stable-query-key-filters	async state	44.0%	0.0%	0.0%	0.0%
02-rn-rq-dependent-query-enabled-gate	async state	36.0%	0.0%	0.0%	0.0%
03-rn-rq-mutation-invalidate-on-success	async state	58.0%	0.0%	0.0%	0.0%
04-rn-rq-optimistic-update-rollback	async state	77.1%	20.0%	77.8%	100.0%
05-rn-zustand-async-action-lifecycle	async state	100.0%	100.0%	100.0%	100.0%
06-rn-zustand-request-id-stale-response-guard	async state	66.0%	0.0%	0.0%	0.0%
07-rn-zustand-persist-async-hydration-gate	async state	55.7%	10.0%	50.0%	100.0%
09-rn-jotai-async-read-atom-suspense	async state	90.0%	70.0%	100.0%	100.0%
10-rn-jotai-loadable-async-atom-no-suspense	async state	25.0%	0.0%	0.0%	0.0%
11-rn-jotai-async-write-atom-mutation	async state	100.0%	100.0%	100.0%	100.0%
12-rn-jotai-atomwithstorage-async-rehydrate-guard	async state	88.3%	40.0%	97.6%	100.0%
13-rn-react-usedeferredvalue-search-stale-indicator	async state	87.5%	60.0%	100.0%	100.0%
14-rn-rq-react-native-appstate-focus-online-managers	async state	33.3%	0.0%	0.0%	0.0%
01-rn-flatlist-scroll-to-unread-fixed-rows	lists	56.0%	0.0%	0.0%	0.0%
02-rn-flatlist-basic-contact-list	lists	100.0%	100.0%	100.0%	100.0%
03-rn-flatlist-two-column-product-grid	lists	100.0%	100.0%	100.0%	100.0%
04-rn-flatlist-viewability-impression-tracking	lists	80.0%	0.0%	0.0%	0.0%
05-rn-flatlist-large-feed-window-tuning	lists	84.0%	30.0%	91.7%	100.0%
06-rn-flatlist-pull-to-refresh-feed	lists	76.0%	0.0%	0.0%	0.0%
07-rn-flatlist-to-sectionlist-conversion	lists	100.0%	100.0%	100.0%	100.0%
08-rn-flashlist-basic-list	lists	93.3%	80.0%	100.0%	100.0%
09-rn-flashlist-horizontal-rail-list-header	lists	72.5%	20.0%	77.8%	100.0%
10-rn-flashlist-horizontal-dynamic-height-header-first-item	lists	78.3%	0.0%	0.0%	0.0%
11-rn-flashlist-expandable-cards-layout-state	lists	36.7%	0.0%	0.0%	0.0%
12-rn-flashlist-chat-start-from-bottom	lists	40.0%	0.0%	0.0%	0.0%
13-rn-flashlist-masonry-photo-gallery	lists	56.0%	0.0%	0.0%	0.0%
14-rn-flatlist-to-legendlist-migration	lists	6.7%	0.0%	0.0%	0.0%
15-rn-legendlist-sticky-grouped-agenda	lists	82.0%	10.0%	50.0%	100.0%
16-rn-legendlist-section-scroll-to-location	lists	94.0%	90.0%	100.0%	100.0%
17-rn-legendlist-chat-start-from-bottom	lists	40.0%	0.0%	0.0%	0.0%
18-rn-legendlist-expandable-cards-recycling-state	lists	37.5%	0.0%	0.0%	0.0%
01-rn-nav-stack-product-details	navigation	72.0%	0.0%	0.0%	0.0%
02-rn-nav-stack-filter-default	navigation	100.0%	100.0%	100.0%	100.0%
03-rn-nav-stack-header-from-param	navigation	97.5%	90.0%	100.0%	100.0%
04-rn-nav-modal-compose-note	navigation	95.0%	90.0%	100.0%	100.0%
05-rn-nav-tabs-three-sections	navigation	96.7%	90.0%	100.0%	100.0%
06-rn-nav-home-tab-nested-details	navigation	83.3%	50.0%	99.6%	100.0%
07-rn-nav-drawer-account-help	navigation	97.5%	90.0%	100.0%	100.0%
09-rn-nav-login-replace-home	navigation	58.0%	0.0%	0.0%	0.0%
10-rn-nav-deeplink-profile-basic	navigation	100.0%	100.0%	100.0%	100.0%
12-rn-nav-unsaved-edit-confirm	navigation	70.0%	0.0%	0.0%	0.0%
13-rn-nav-focus-refresh-inbox	navigation	100.0%	100.0%	100.0%	100.0%
14-rn-nav-tabs-lazy-first-render	navigation	100.0%	100.0%	100.0%	100.0%
15-rn-nav-header-action-stateful	navigation	80.0%	80.0%	100.0%	100.0%
01-rn-stylesheet-card-variants	react native apis	60.0%	40.0%	97.6%	100.0%
02-rn-stylesheet-absolute-fill-overlay-badge	react native apis	63.3%	0.0%	0.0%	0.0%
03-rn-stylesheet-divider-hairline	react native apis	76.7%	30.0%	91.7%	100.0%
04-rn-image-network-dimensions	react native apis	96.7%	90.0%	100.0%	100.0%
05-rn-imagebackground-resize-mode	react native apis	97.5%	90.0%	100.0%	100.0%
06-rn-platform-select-small-style-delta	react native apis	100.0%	100.0%	100.0%	100.0%
07-rn-platform-os-conditional-cta	react native apis	100.0%	100.0%	100.0%	100.0%
08-rn-platform-file-extension-component-split	react native apis	100.0%	100.0%	100.0%	100.0%
09-rn-keyboardavoidingview-login-form	react native apis	100.0%	100.0%	100.0%	100.0%

Table 55: Per-eval results (score and pass@k) for minimax-m2.7



Figure 35: Benchmark Plot for minimax-m2.7.

vercel/minimax/minimax-m2.7 · Run → Category → Pass/Fail




Figure 36: Sankey per run for minimax-m2.7 showing the distribution of evals with respect to categories and results

A.18 qwen2.5-coder-32b-instruct

Model	Runs	Score	Reqs	P@1	P@5	P@10
cloudflare-ai-gateway/workers-ai/@cf/qwen/qwen2.5-	10	51.2%	1316/2570	18.8%	29.3%	28.8%

Table 56: Model summary for qwen2.5-coder-32b-instruct

Category	Score	Evals	Requirements	Failed requirements
animation	25.9%	122	53	389
async state	44.4%	121	65	329
lists	57.3%	176	79	318
navigation	66.1%	111	45	140
react native apis	70.2%	86	31	78

Table 57: Results by category for qwen2.5-coder-32b-instruct. Score: mean(per-eval passed/total weight) across evals and runs. Evals: count of evaluation tasks. Requirements: sum of requirement count per eval. Failed requirements: sum of failed requirement instances across runs.

Eval	Category	Score	P@1	P@5	P@10
01-rn-anim-pressable-scale-with-timing	animation	53.3%	0.0%	0.0%	0.0%
02-rn-anim-spring-toggle-card-height	animation	40.7%	0.0%	0.0%	0.0%
03-rn-anim-layout-enter-exit-list-items	animation	90.0%	80.0%	100.0%	100.0%
04-rn-anim-scroll-linked-collapsing-header	animation	46.7%	20.0%	77.8%	100.0%
05-rn-anim-pan-drag-card-with-snap-points	animation	10.0%	0.0%	0.0%	0.0%
06-rn-anim-longpress-then-pan-activation-gate	animation	22.5%	0.0%	0.0%	0.0%
07-rn-anim-pinch-pan-simultaneous-photo-canvas	animation	48.0%	0.0%	0.0%	0.0%
09-rn-anim-android-modal-gesture-root-safety	animation	3.7%	0.0%	0.0%	0.0%
10-rn-anim-worklet-thread-bridge-schedule-on-rn	animation	17.5%	0.0%	0.0%	0.0%
11-rn-rngh-exclusive-single-double-tap	animation	2.0%	0.0%	0.0%	0.0%
12-rn-rngh-scroll-pan-failure-chain	animation	4.4%	0.0%	0.0%	0.0%
13-rn-worklets-runtime-offload-and-bridge	animation	6.7%	0.0%	0.0%	0.0%
14-rn-keyboard-controller-aware-scroll-footer	animation	2.0%	0.0%	0.0%	0.0%
01-rn-rq-stable-query-key-filters	async state	10.0%	0.0%	0.0%	0.0%
02-rn-rq-dependent-query-enabled-gate	async state	18.0%	0.0%	0.0%	0.0%
03-rn-rq-mutation-invalidate-on-success	async state	22.0%	0.0%	0.0%	0.0%
04-rn-rq-optimistic-update-rollback	async state	62.9%	0.0%	0.0%	0.0%
05-rn-zustand-async-action-lifecycle	async state	92.0%	60.0%	100.0%	100.0%
06-rn-zustand-request-id-stale-response-guard	async state	22.0%	0.0%	0.0%	0.0%
07-rn-zustand-persist-async-hydration-gate	async state	41.4%	0.0%	0.0%	0.0%
09-rn-jotai-async-read-atom-suspense	async state	22.5%	0.0%	0.0%	0.0%
10-rn-jotai-loadable-async-atom-no-suspense	async state	12.5%	0.0%	0.0%	0.0%
11-rn-jotai-async-write-atom-mutation	async state	66.0%	0.0%	0.0%	0.0%
12-rn-jotai-atomwithstorage-async-rehydrate-guard	async state	60.0%	0.0%	0.0%	0.0%
13-rn-react-usedeferredvalue-search-stale-indicator	async state	91.7%	66.7%	100.0%	66.7%
14-rn-rq-react-native-appstate-focus-online-managers	async state	33.3%	0.0%	0.0%	0.0%
01-rn-flatlist-scroll-to-unread-fixed-rows	lists	42.0%	0.0%	0.0%	0.0%
02-rn-flatlist-basic-contact-list	lists	66.7%	0.0%	0.0%	0.0%
03-rn-flatlist-two-column-product-grid	lists	80.0%	40.0%	97.6%	100.0%
04-rn-flatlist-viewability-impression-tracking	lists	66.0%	0.0%	0.0%	0.0%
05-rn-flatlist-large-feed-window-tuning	lists	90.0%	50.0%	99.6%	100.0%
06-rn-flatlist-pull-to-refresh-feed	lists	84.0%	20.0%	77.8%	100.0%
07-rn-flatlist-to-sectionlist-conversion	lists	80.0%	0.0%	0.0%	0.0%
08-rn-flashlist-basic-list	lists	66.7%	0.0%	0.0%	0.0%
09-rn-flashlist-horizontal-rail-list-header	lists	75.0%	20.0%	77.8%	100.0%
10-rn-flashlist-horizontal-dynamic-height-header-first-item	lists	75.0%	0.0%	0.0%	0.0%
11-rn-flashlist-expandable-cards-layout-state	lists	33.3%	0.0%	0.0%	0.0%
12-rn-flashlist-chat-start-from-bottom	lists	40.0%	0.0%	0.0%	0.0%
13-rn-flashlist-masonry-photo-gallery	lists	77.1%	0.0%	0.0%	0.0%
14-rn-flatlist-to-legendlist-migration	lists	0.0%	0.0%	0.0%	0.0%
15-rn-legendlist-sticky-grouped-agenda	lists	46.0%	0.0%	0.0%	0.0%
16-rn-legendlist-section-scroll-to-location	lists	62.0%	30.0%	91.7%	100.0%
17-rn-legendlist-chat-start-from-bottom	lists	30.0%	0.0%	0.0%	0.0%
18-rn-legendlist-expandable-cards-recycling-state	lists	25.0%	0.0%	0.0%	0.0%
01-rn-nav-stack-product-details	navigation	60.0%	0.0%	0.0%	0.0%
02-rn-nav-stack-filter-default	navigation	94.0%	80.0%	100.0%	100.0%
03-rn-nav-stack-header-from-param	navigation	56.2%	12.5%	62.5%	12.5%
04-rn-nav-modal-compose-note	navigation	75.0%	0.0%	0.0%	0.0%
05-rn-nav-tabs-three-sections	navigation	73.3%	20.0%	100.0%	20.0%
06-rn-nav-home-tab-nested-details	navigation	43.3%	10.0%	50.0%	100.0%
07-rn-nav-drawer-account-help	navigation	50.0%	0.0%	0.0%	0.0%
09-rn-nav-login-replace-home	navigation	42.0%	0.0%	0.0%	0.0%
10-rn-nav-deeplink-profile-basic	navigation	100.0%	100.0%	100.0%	100.0%
12-rn-nav-unsaved-edit-confirm	navigation	55.0%	0.0%	0.0%	0.0%
13-rn-nav-focus-refresh-inbox	navigation	100.0%	100.0%	100.0%	100.0%
14-rn-nav-tabs-lazy-first-render	navigation	66.7%	0.0%	0.0%	0.0%
15-rn-nav-header-action-stateful	navigation	44.4%	0.0%	0.0%	0.0%
01-rn-stylesheet-card-variants	react native apis	73.3%	60.0%	100.0%	100.0%
02-rn-stylesheet-absolute-fill-overlay-badge	react native apis	30.0%	0.0%	0.0%	0.0%
03-rn-stylesheet-divider-hairline	react native apis	16.7%	0.0%	0.0%	0.0%
04-rn-image-network-dimensions	react native apis	33.3%	0.0%	0.0%	0.0%
05-rn-imagebackground-resize-mode	react native apis	90.0%	70.0%	100.0%	100.0%
06-rn-platform-select-small-style-delta	react native apis	100.0%	100.0%	100.0%	100.0%
07-rn-platform-os-conditional-cta	react native apis	100.0%	100.0%	100.0%	100.0%
08-rn-platform-file-extension-component-split	react native apis	100.0%	100.0%	100.0%	100.0%
09-rn-keyboardavoidingview-login-form	react native apis	100.0%	100.0%	100.0%	100.0%

Table 58: Per-eval results (score and pass@k) for qwen2.5-coder-32b-instruct

Benchmark (10 runs) · cloudflare-ai-gateway/workers-ai/@cf/qwen/qwen2.5-coder-32b-instruct · score 51.2% (1316/2570 reqs) · pass@1/5/10: 18.8% / 29.3% / 28.8% · cost \$4.98



Figure 37: Benchmark Plot for qwen2.5-coder-32b-instruct.

cloudflare-ai-gateway/workers-ai/@cf/qwen/qwen2.5-coder-32b-instruct · Run → Category →




Figure 38: Sankey per run for qwen2.5-coder-32b-instruct showing the distribution of evals with respect to categories and results